



# Project 2

[https://github.com/DX0ZART/Group\\_3](https://github.com/DX0ZART/Group_3)

## Load Libraries

```
In [ ]: import numpy as np
import pandas as pd
from matplotlib import pyplot as plt
import tensorflow as tf
import os

from sklearn.model_selection import train_test_split

from tensorflow.python.keras.utils import np_utils
from tensorflow.keras.models import Sequential, Model

from sklearn.metrics import classification_report, confusion_matrix, accuracy_

from tensorflow.keras.layers import Dense, Dropout, Flatten, Activation, Batch
from tensorflow.keras.layers import Conv2D, MaxPooling2D

from tensorflow.keras.callbacks import EarlyStopping, ReduceLROnPlateau, Model
from tensorflow.keras.applications import VGG19, ResNet50, InceptionV3, Efficie
# from tensorflow.python.keras.layers.convolutional import Conv2D, MaxPooling2
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.optimizers import Adam, SGD, Adagrad, Adadelta, RMSprop
```

## 0. Data Loading and Preprocessing

### Load Data

```
In [ ]: # Extracting all filenames iteratively
base_path = 'COVID-19_Radiography_Dataset'
categories = ['COVID/images', 'Normal/images', 'Viral Pneumonia/images']

# load file names to fnames list object
fnames = []
for category in categories:
    image_folder = os.path.join(base_path, category)
    file_names = os.listdir(image_folder)
    full_path = [os.path.join(image_folder, file_name) for file_name in file_n
fnames.append(full_path)

print('number of images for each category:', [len(f) for f in fnames])
# print(fnames[0:2]) #examples of file names
```

number of images for each category: [3616, 10192, 1345]

```
In [ ]: # Reduce number of images to first 1345 for each category  
# fnames[0]=fnames[0][0:1344]  
# fnames[1]=fnames[1][0:1344]  
# fnames[2]=fnames[2][0:1344]
```

## Preprocessing

```
In [ ]: # Import image, load to array of shape height, width, channels, then min/max t  
# Write preprocessor that will match up with model's expected input shape.  
from keras.preprocessing import image  
import numpy as np  
from PIL import Image  
  
def preprocessor(img_path):  
    img = Image.open(img_path).convert("RGB").resize((192,192)) # import image  
    img = (np.float32(img)-1.)/(255-1.) # min max transformation  
    img=img.reshape((192,192,3)) # Create final shape as array with correct di  
    return img  
  
#Try on single flower file (imports file and preprocesses it to data with foll  
preprocessor('COVID-19_Radiography_Dataset/COVID/images/COVID-2273.png').shape
```

```
Out[ ]: (192, 192, 3)
```

```
In [ ]: #Import image files iteratively and preprocess them into array of data  
  
# aggregate filepath  
image_filepaths=fnames[0]+fnames[1]+fnames[2]  
  
# map functions apply your preprocessor function one step at a time to each fi  
preprocessed_image_data=list(map(preprocessor,image_filepaths ))  
  
# cast the list into array for Keras  
X= np.array(preprocessed_image_data)
```

```
In [ ]: len(image_filepaths)
```

```
Out[ ]: 15153
```

```
In [ ]: print(len(X) ) #double check the numbers if the files  
print(X.shape ) #making sure the dimensions for all images are correct  
print(X.min().round() ) #min value of every image is zero  
print(X.max() ) #max value of every image is one
```

```
15153  
(15153, 192, 192, 3)  
-0.0  
1.0
```

```
In [ ]: len(fnames[2])
```

```
Out[ ]: 1345
```

```
In [ ]: # Create y data made up of correctly ordered labels from file folders
from itertools import repeat

#...corresponding to each flower type

print('number of images for each category:', [len(f) for f in fnames])
covid=list(repeat("COVID", len(fnames[0])))
normal=list(repeat("NORMAL", len(fnames[1])))
pneumonia=list(repeat("PNEUMONIA", len(fnames[2])))

#combine into single list of y labels
y_labels = covid+normal+pneumonia

#check length, same as X above
print(len(y_labels) )

# Need to one hot encode for Keras.

import pandas as pd
y=pd.get_dummies(y_labels)

display(y)
```

number of images for each category: [3616, 10192, 1345]  
15153

	COVID	NORMAL	PNEUMONIA
<b>0</b>	True	False	False
<b>1</b>	True	False	False
<b>2</b>	True	False	False
<b>3</b>	True	False	False
<b>4</b>	True	False	False
...	...	...	...
<b>15148</b>	False	False	True
<b>15149</b>	False	False	True
<b>15150</b>	False	False	True
<b>15151</b>	False	False	True
<b>15152</b>	False	False	True

15153 rows × 3 columns

```
In [ ]: import matplotlib.pyplot as plt
from mpl_toolkits.axes_grid1 import ImageGrid
import numpy as np

im1 =preprocessor(fnames[0][0])
```

```

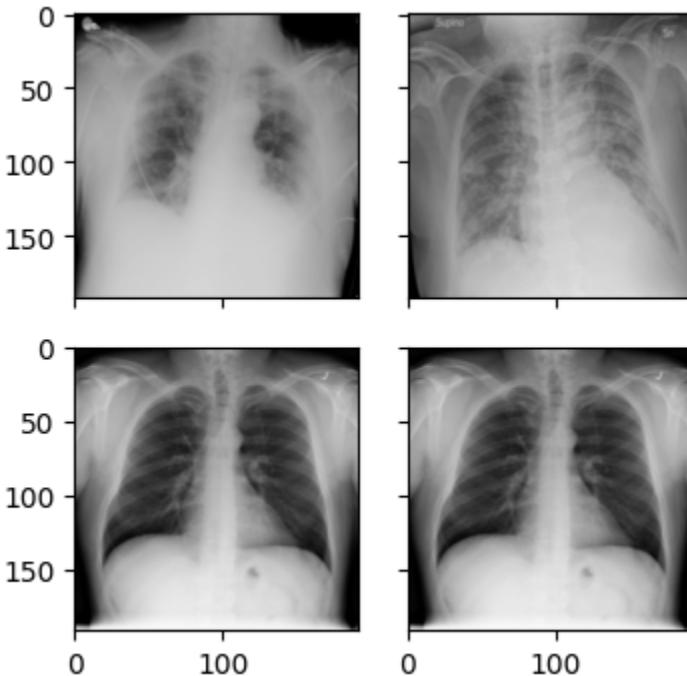
im2 =preprocessor(fnames[0][1])
im3 =preprocessor(fnames[1][1])
im4 =preprocessor(fnames[1][1])

fig = plt.figure(figsize=(4., 4.))
grid = ImageGrid(fig, 111,
                  nrows_ncols=(2, 2), # creates 2x2 grid of axes
                  axes_pad=0.25, # pad between axes in inch.
                  )

for ax, im in zip(grid, [im1, im2, im3, im4]):
    ax.imshow(im)
plt.show()

```

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers). Got range [-0.003937008..0.8425197].  
 Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers). Got range [-0.003937008..0.8897638].  
 Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers). Got range [-0.003937008..0.98031497].  
 Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers). Got range [-0.003937008..0.98031497].



## Train Test Split

```

In [ ]: # =====Train test split resized images (Hackathon Note!! Use same train test
from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y, stratify = y, test_s
X_test.shape, y_test.shape

```

```
Out[ ]: ((4849, 192, 192, 3), (4849, 3))
```

```
In [ ]: #Clear objects from memory
del(X)
del(y)
del(preprocessed_image_data)
```

## Save to Files

```
In [ ]: #Save data to be able to reload quickly if crashes
import pickle

with open('data/X_train.pkl', 'wb') as file:
    # A new file will be created
    pickle.dump(X_train, file)

with open('data/X_test.pkl', 'wb') as file:
    pickle.dump(X_test, file)

with open('data/y_train.pkl', 'wb') as file:
    pickle.dump(y_train, file)

with open('data/y_test.pkl', 'wb') as file:
    pickle.dump(y_test, file)
```

## 1. Dataset and Exploratory Data Analysis

### Load Data

```
In [ ]: import pickle

# Open the file in binary mode
with open('data/X_train.pkl', 'rb') as file:
    # Call load method to deserialize
    X_train = pickle.load(file)

# Open the file in binary mode
with open('data/y_train.pkl', 'rb') as file:
    # Call load method to deserialize
    y_train = pickle.load(file)

# Open the file in binary mode
with open('data/X_test.pkl', 'rb') as file:
    # Call load method to deserialize
    X_test = pickle.load(file)

# Open the file in binary mode
with open('data/y_test.pkl', 'rb') as file:
    # Call load method to deserialize
    y_test = pickle.load(file)
```

## Describe the dataset - basic statistics

```
In [ ]: import seaborn as sns
```

```
        sns.set_style('whitegrid')
        plt.rcParams['figure.figsize'] = (12, 8)
```

```
In [ ]: # Dataset Overview
print("Dataset Description:")
print("This dataset contains chest X-ray images for COVID-19 classification.")
print("The images are categorized into three classes:")
print("  1. COVID: X-rays showing COVID-19 positive cases")
print("  2. NORMAL: X-rays from healthy individuals")
print("  3. PNEUMONIA: X-rays showing viral pneumonia (non-COVID)")
print("==" * 80)
# Class Labels
print("CLASS LABELS:")
print(f"Number of classes: {y_train.shape[1]}")
print(f"Class names: {list(y_train.columns)}")
```

Dataset Description:

This dataset contains chest X-ray images for COVID-19 classification.

The images are categorized into three classes:

1. COVID: X-rays showing COVID-19 positive cases
  2. NORMAL: X-rays from healthy individuals
  3. PNEUMONIA: X-rays showing viral pneumonia (non-COVID)
- 

=

CLASS LABELS:

Number of classes: 3

Class names: ['COVID', 'NORMAL', 'PNEUMONIA']

```
In [ ]: # Data Dimensions
```

```
print("DATA DIMENSIONS:")
print("Training Set:")
print(f"  - X_train shape: {X_train.shape}")
print(f"  - y_train shape: {y_train.shape}")
print(f"  - Number of training samples: {X_train.shape[0]},")
print("Test Set:")
print(f"  - X_test shape: {X_test.shape}")
print(f"  - y_test shape: {y_test.shape}")
print(f"  - Number of test samples: {X_test.shape[0]},")
print(f"Total Dataset Size: {X_train.shape[0] + X_test.shape[0]}, images")
print("Image Properties:")
print(f"  - Image dimensions: {X_train.shape[1]} x {X_train.shape[2]} pixels")
print(f"  - Color channels: {X_train.shape[3]} (RGB)")
print(f"  - Pixel value range: [{X_train.min():.3f}, {X_train.max():.3f}] (nor
```

DATA DIMENSIONS:  
Training Set:  
- X\_train shape: (10304, 192, 192, 3)  
- y\_train shape: (10304, 3)  
- Number of training samples: 10,304  
Test Set:  
- X\_test shape: (4849, 192, 192, 3)  
- y\_test shape: (4849, 3)  
- Number of test samples: 4,849  
Total Dataset Size: 15,153 images  
Image Properties:  
- Image dimensions: 192 x 192 pixels  
- Color channels: 3 (RGB)  
- Pixel value range: [-0.004, 1.000] (normalized)

## Basic Statistics - Image Samples

```
In [ ]: def display_class_samples(X, y, n_samples=5):
    classes = y.columns.tolist()
    n_classes = len(classes)

    fig, axes = plt.subplots(n_classes, n_samples, figsize=(n_samples * 3, n_c
fig.suptitle('Sample X-Ray Images from Each Class', fontsize=16, fontweight='bold')

    for i, class_name in enumerate(classes):
        class_indices = np.where(y[class_name] == 1)[0]
        sample_indices = np.random.choice(class_indices, n_samples, replace=False)

        for j, idx in enumerate(sample_indices):
            ax = axes[i, j] if n_classes > 1 else axes[j]
            ax.imshow(X[idx])
            ax.axis('off')

            if j == 0:
                ax.set_title(f'{class_name}\nSample {j+1}', fontweight='bold',
            else:
                ax.set_title(f'Sample {j+1}', fontsize=10)

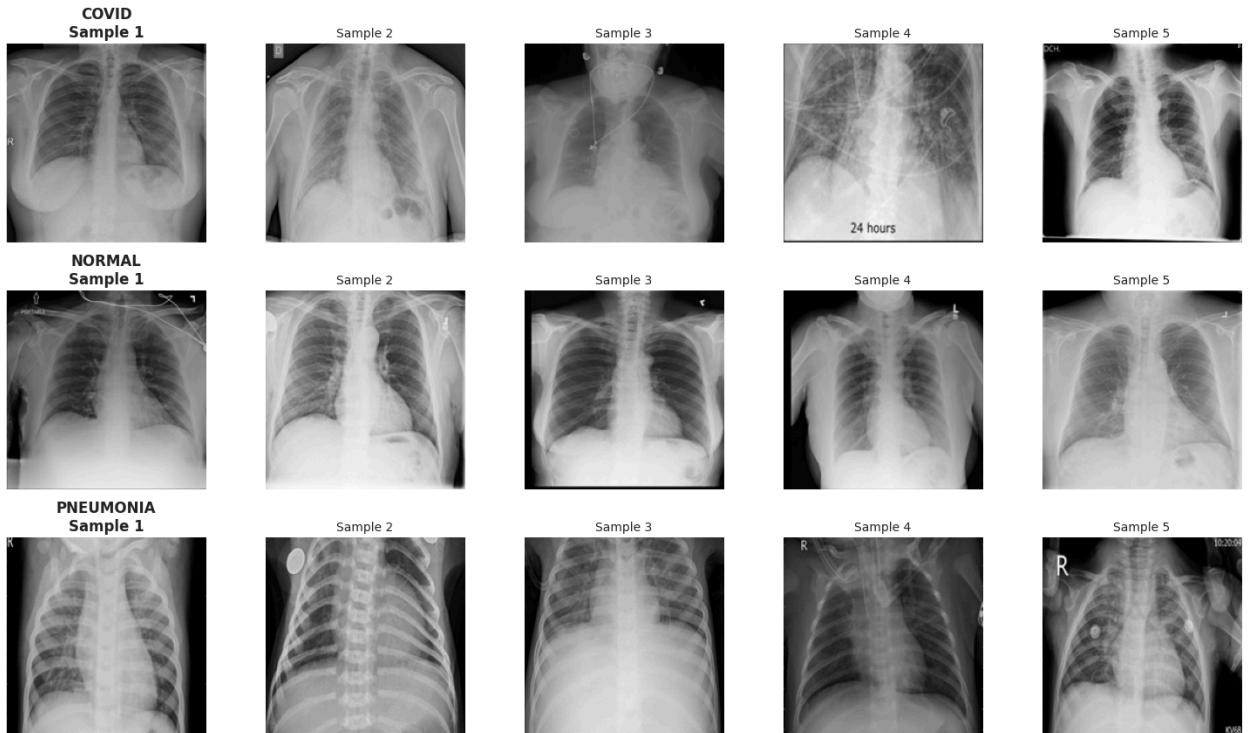
    plt.tight_layout()
    plt.show()

print("Sample images from the training set:")
display_class_samples(X_train, y_train, n_samples=5)
```

```
Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers). Got range [-0.003937008..0.8307087].  
Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers). Got range [-0.003937008..0.8543307].  
Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers). Got range [-0.003937008..1.0].  
Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers). Got range [-0.003937008..0.9488189].  
Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers). Got range [-0.003937008..0.96062994].  
Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers). Got range [-0.003937008..0.96456695].  
Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers). Got range [-0.003937008..0.88188976].  
Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers). Got range [-0.003937008..1.0].  
Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers). Got range [-0.003937008..0.8543307].  
Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers). Got range [-0.003937008..0.9094488].  
Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers). Got range [-0.003937008..0.86220473].  
Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers). Got range [-0.003937008..1.0].
```

Sample images from the training set:

Sample X-Ray Images from Each Class



## Check Class Balance

```
In [ ]: # Calculate class distribution  
train_class_counts = y_train.sum().to_dict()
```

```

test_class_counts = y_test.sum().to_dict()
total_class_counts = {k: train_class_counts[k] + test_class_counts[k] for k in total_class_counts}

# Create DataFrame for better visualization
distribution_df = pd.DataFrame({
    'Training': train_class_counts,
    'Test': test_class_counts,
    'Total': total_class_counts
})

# Calculate percentages
distribution_df['Train %'] = (distribution_df['Training'] / distribution_df['Total']) * 100
distribution_df['Test %'] = (distribution_df['Test'] / distribution_df['Total']) * 100
distribution_df['Total %'] = (distribution_df['Total'] / distribution_df['Total']) * 100

print("Class Distribution:")
print("-" * 80)
print(distribution_df)

print("\nDATASET IS NOT BALANCED")
print("Balancing Required")

```

Class Distribution:

	Training	Test	Total	Train %	Test %	Total %
COVID	2459	1157	3616	23.86	23.86	23.86
NORMAL	6930	3262	10192	67.26	67.27	67.26
PNEUMONIA	915	430	1345	8.88	8.87	8.88

DATASET IS NOT BALANCED

Balancing Required

## Strategies to Handle Class Imbalance

1. **Class Weighting:** Assign higher weights to minority classes during training. It does not change the dataset size and penalizes misclassifications of minority classes more heavily.
2. **Oversampling:** Increase minority class samples (e.g., Randomly duplicate samples from minority classes to match majority class.)
3. **Undersampling:** Randomly remove samples from majority classes to match minority class.
4. **Augmentation:** Generate synthetic samples through transformations
  - I would choose **this method**.

## Data Augmentation:

- Creates diverse synthetic samples without exact duplicates
- Improves model generalization and robustness
- Particularly effective for image data
- Reduces overfitting risk

### Expected Performance Impact:

- More balanced precision/recall across classes
- Better sensitivity (recall) for minority classes
- Slight decrease in overall accuracy but significant improvement in per-class metrics
- Higher F1-scores for underrepresented classes

### Strategy: Data Augmentation

```
In [ ]: idx_covid = np.where(y_train["COVID"] == 1)[0]
idx_normal = np.where(y_train["NORMAL"] == 1)[0]
idx_pneu = np.where(y_train["PNEUMONIA"] == 1)[0]

X_covid, X_normal, X_pneu = X_train[idx_covid], X_train[idx_normal], X_train[i
y_covid, y_normal, y_pneu = y_train.iloc[idx_covid], y_train.iloc[idx_normal],
```

```
In [ ]: datagen = ImageDataGenerator(
    rotation_range=15, # Rotation: ±15 degrees
    width_shift_range=0.1, # Width shift: ±10%
    height_shift_range=0.1, # Height shift: ±10%
    horizontal_flip=True, # Horizontal flip: Enabled
    zoom_range=0.1, # Zoom: ±10%
    fill_mode='nearest'
)
```

```
In [ ]: def augment_class(X_class, y_class, target_size):
    n_to_generate = target_size - len(X_class)
    if n_to_generate <= 0:
        return X_class, y_class

    aug_iter = datagen.flow(X_class, y_class, batch_size=32)
    X_aug, y_aug = [], []

    for _ in range(n_to_generate // 32 + 1):
        x_batch, y_batch = next(aug_iter)
        X_aug.append(x_batch)
        y_aug.append(y_batch)
        if len(np.concatenate(X_aug)) >= target_size - len(X_class):
            break

    X_aug = np.concatenate(X_aug)[:n_to_generate]
    y_aug = np.concatenate(y_aug)[:n_to_generate]
    return np.concatenate([X_class, X_aug]), np.concatenate([y_class, y_aug])
```

```
In [ ]: target = len(X_normal) # 6930  
  
X_covid_aug, y_covid_aug = augment_class(X_covid, y_covid, target)  
X_pneu_aug, y_pneu_aug = augment_class(X_pneu, y_pneu, target)
```

```
In [ ]: # In case OOM  
del(X_test)  
del(y_train)
```

```
In [ ]: X_balanced = np.concatenate([X_normal, X_covid_aug, X_pneu_aug])  
y_balanced = np.concatenate([y_normal, y_covid_aug, y_pneu_aug])  
  
p = np.random.permutation(len(X_balanced))  
X_train_balanced, y_train_balanced = X_balanced[p], y_balanced[p]
```

```
In [ ]: # In case OOM  
del(X_covid_aug)  
del(y_covid_aug)  
del(X_balanced)  
del(y_balanced)  
del(p)
```

## Check Balance Again

```
In [ ]: y_train_balanced = pd.DataFrame(y_train_balanced, columns = ["COVID", "NORMAL"]  
  
In [ ]: # Calculate class distribution  
train_class_counts = y_train_balanced.sum().to_dict()  
test_class_counts = y_test.sum().to_dict()  
total_class_counts = {k: train_class_counts[k] + test_class_counts[k] for k in train_class_counts}  
  
# Create DataFrame for better visualization  
distribution_df = pd.DataFrame({  
    'Training': train_class_counts,  
    'Test': test_class_counts,  
    'Total': total_class_counts  
})  
  
# Calculate percentages  
distribution_df['Train %'] = (distribution_df['Training'] / distribution_df['Total']) * 100  
distribution_df['Test %'] = (distribution_df['Test'] / distribution_df['Total']) * 100  
distribution_df['Total %'] = (distribution_df['Total'] / distribution_df['Total']) * 100  
  
print("Class Distribution:")  
print("-" * 80)  
print(distribution_df)  
  
print("\nDATASET IS NOW BALANCED")
```

Class Distribution:

	Training	Test	Total	Train %	Test %	Total %
COVID	6930	1157	8087	33.46	23.86	31.64
NORMAL	6930	3262	10192	33.46	67.27	39.87
PNEUMONIA	6853	430	7283	33.09	8.87	28.49

DATASET IS NOW BALANCED

```
In [ ]: # Visualize augmented samples
print("Example of augmented images:")
print("(Original image followed by augmented versions)\n")

sample_image = X_train[0:1] # Keep batch dimension

fig, axes = plt.subplots(2, 3, figsize=(15, 10))
axes = axes.ravel()

# Show original
axes[0].imshow(sample_image[0])
axes[0].set_title('Original Image', fontweight='bold', fontsize=12)
axes[0].axis('off')

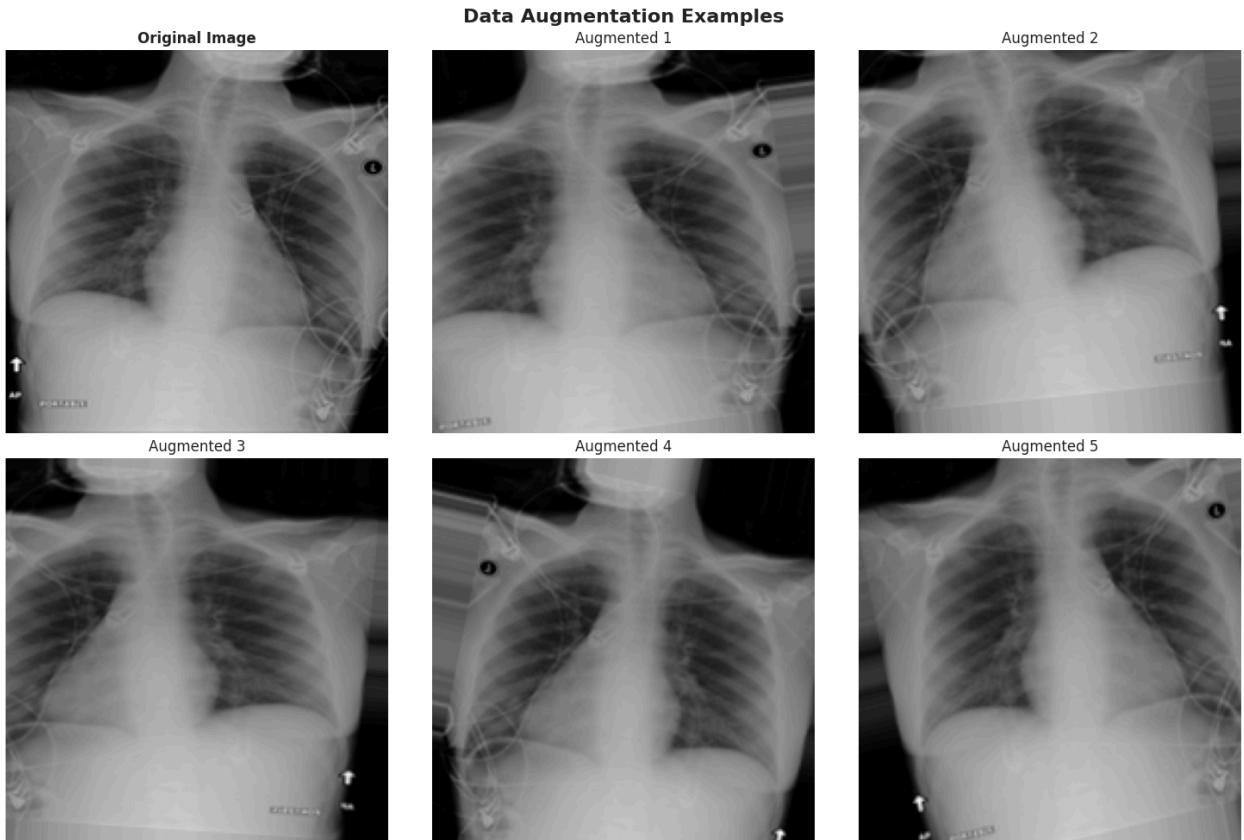
# Generate and show augmented versions
datagen.fit(sample_image)
for i, batch in enumerate(datagen.flow(sample_image, batch_size=1)):
    if i >= 5:
        break
    axes[i+1].imshow(batch[0])
    axes[i+1].set_title(f'Augmented {i+1}', fontsize=12)
    axes[i+1].axis('off')

plt.suptitle('Data Augmentation Examples', fontsize=16, fontweight='bold', y=0.95)
plt.tight_layout()
plt.show()
```

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers). Got range [-0.003937008..1.0].  
Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers). Got range [-0.003937008..0.8098158].  
Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers). Got range [-0.0038251434..0.9629155].  
Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers). Got range [-0.0038858745..0.9592573].  
Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers). Got range [-0.003937008..0.9869442].  
Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers). Got range [-0.003937008..0.9843035].

Example of augmented images:

(Original image followed by augmented versions)



## Save to Files

```
In [ ]: # In case OOM
del(X_train)
del(y_test)
```

```
In [ ]: #Save data to be able to reload quickly if memory crashes or if you run RuntimeError
import pickle

# Open a file and use dump()
with open('data/X_train_balanced.pkl', 'wb') as file:
    # A new file will be created
    pickle.dump(X_train_balanced, file)

with open('data/y_train_balanced.pkl', 'wb') as file:
    # A new file will be created
    pickle.dump(y_train_balanced, file)
```

## Real-World Impact and Beneficiaries

This COVID-19 X-ray classification model has significant practical applications:

### 1. Healthcare Providers:

- **Emergency Departments:** Rapid triage of patients with respiratory

- symptoms
- **Radiologists:** AI-assisted diagnosis to reduce workload and catch missed cases
- **Telemedicine:** Remote diagnosis support for areas without specialist access

## 2. Public Health:

- **Outbreak Monitoring:** Early detection of COVID-19 clusters in communities
- **Resource Allocation:** Prioritizing limited RT-PCR tests for high-risk cases

## 3. Patients:

- **Faster Diagnosis:** Reduced wait time for results (minutes vs. hours/days)
- **Early Treatment:** Quicker initiation of appropriate care

## 2. Baseline CNN Model

### Clean Memory

```
In [ ]: import gc  
gc.collect()
```

```
Out[ ]: 0
```

```
In [ ]: # Load preprocessed data  
import pickle  
  
with open('data/X_train.pkl', 'rb') as file:  
    X_train = pickle.load(file)  
  
with open('data/y_train.pkl', 'rb') as file:  
    y_train = pickle.load(file)  
  
with open('data/X_test.pkl', 'rb') as file:  
    X_test = pickle.load(file)  
  
with open('data/y_test.pkl', 'rb') as file:  
    y_test = pickle.load(file)  
  
print(f"Training data shape: {X_train.shape}")  
print(f"Training labels shape: {y_train.shape}")  
print(f"Test data shape: {X_test.shape}")  
print(f"Test labels shape: {y_test.shape}")
```

```
Training data shape: (10304, 192, 192, 3)
Training labels shape: (10304, 3)
Test data shape: (4849, 192, 192, 3)
Test labels shape: (4849, 3)
```

```
In [ ]: # Create validation split from training data
X_train_split, X_val, y_train_split, y_val = train_test_split(
    X_train, y_train,
    test_size=0.15,
    random_state=42,
    stratify=y_train
)

print(f"Training split shape: {X_train_split.shape}")
print(f"Validation split shape: {X_val.shape}")
print(f"Test split shape: {X_test.shape}")
```

```
Training split shape: (8758, 192, 192, 3)
Validation split shape: (1546, 192, 192, 3)
Test split shape: (4849, 192, 192, 3)
```

## Baseline: basic Convolutional Neural Network

Architecture:

- 3 Convolutional blocks (Conv2D → BatchNorm → MaxPooling → Dropout)
- Dense layers with dropout for classification
- Output layer with softmax for 3-class classification

Configuration:

- Loss Function: Categorical Crossentropy (multi-class classification)
- Optimizer: Adam (learning\_rate=0.001)
- Metrics: Accuracy, Precision, Recall
- Batch Size: 32
- Epochs: 30 with early stopping

```
In [ ]: def create_baseline_cnn(input_shape=(192, 192, 3), num_classes=3):
    """
    Create a baseline CNN model.

    Architecture:
    - Conv Block 1: 32 filters, 3x3 kernel
    - Conv Block 2: 64 filters, 3x3 kernel
    - Conv Block 3: 128 filters, 3x3 kernel
    - Dense Layer: 256 units
    - Output Layer: 3 units (softmax)
    """
    model = Sequential([
        # First Conv Block
        Conv2D(32, (3, 3), activation='relu', padding='same', input_shape=input_shape),
        BatchNormalization(),
        MaxPooling2D(),
        Dropout(0.2),
        # Second Conv Block
        Conv2D(64, (3, 3), activation='relu', padding='same'),
        BatchNormalization(),
        MaxPooling2D(),
        Dropout(0.2),
        # Third Conv Block
        Conv2D(128, (3, 3), activation='relu', padding='same'),
        BatchNormalization(),
        MaxPooling2D(),
        Dropout(0.2),
        # Dense Layer
        Dense(256, activation='relu'),
        BatchNormalization(),
        Dropout(0.5),
        # Output Layer
        Dense(num_classes, activation='softmax')
    ])
    return model
```

```

        BatchNormalization(),
        Conv2D(32, (3, 3), activation='relu', padding='same'),
        BatchNormalization(),
        MaxPooling2D(pool_size=(2, 2)),
        Dropout(0.25),

        # Second Conv Block
        Conv2D(64, (3, 3), activation='relu', padding='same'),
        BatchNormalization(),
        Conv2D(64, (3, 3), activation='relu', padding='same'),
        BatchNormalization(),
        MaxPooling2D(pool_size=(2, 2)),
        Dropout(0.25),

        # Third Conv Block
        Conv2D(128, (3, 3), activation='relu', padding='same'),
        BatchNormalization(),
        Conv2D(128, (3, 3), activation='relu', padding='same'),
        BatchNormalization(),
        MaxPooling2D(pool_size=(2, 2)),
        Dropout(0.25),

        # Dense Layers
        Flatten(),
        Dense(256, activation='relu'),
        BatchNormalization(),
        Dropout(0.5),
        Dense(128, activation='relu'),
        BatchNormalization(),
        Dropout(0.5),

        # Output Layer
        Dense(num_classes, activation='softmax')
    ], name='Baseline_CNN')

    return model

```

```

In [ ]: # Create and compile the baseline model
baseline_model = create_baseline_cnn()

baseline_model.compile(
    optimizer=Adam(learning_rate=0.001),
    loss='categorical_crossentropy',
    metrics=['accuracy',
              tf.keras.metrics.Precision(name='precision'),
              tf.keras.metrics.Recall(name='recall')])
)

baseline_model.summary()

```

```
/mnt/f/BaiduSyncdisk/1_Classes/3_25Fall/1_ADS/HW/HW_3A_Due_2025.10.12/.venv/lib/python3.12/site-packages/keras/src/layers/convolutional/base_conv.py:113: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.  
    super().__init__(activity_regularizer=activity_regularizer, **kwargs)  
WARNING: All log messages before absl::InitializeLog() is called are written to STDERR  
I0000 00:00:1761860552.692695 32050 gpu_device.cc:2020] Created device /job:localhost/replica:0/task:0/device:GPU:0 with 21458 MB memory: -> device: 0, name: NVIDIA GeForce RTX 4090, pci bus id: 0000:01:00.0, compute capability: 8.9  
Model: "Baseline_CNN"
```

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 192, 192, 32)	896
batch_normalization (BatchNormalization)	(None, 192, 192, 32)	128
conv2d_1 (Conv2D)	(None, 192, 192, 32)	9,248
batch_normalization_1 (BatchNormalization)	(None, 192, 192, 32)	128
max_pooling2d (MaxPooling2D)	(None, 96, 96, 32)	0
dropout (Dropout)	(None, 96, 96, 32)	0
conv2d_2 (Conv2D)	(None, 96, 96, 64)	18,496
batch_normalization_2 (BatchNormalization)	(None, 96, 96, 64)	256
conv2d_3 (Conv2D)	(None, 96, 96, 64)	36,928
batch_normalization_3 (BatchNormalization)	(None, 96, 96, 64)	256
max_pooling2d_1 (MaxPooling2D)	(None, 48, 48, 64)	0
dropout_1 (Dropout)	(None, 48, 48, 64)	0
conv2d_4 (Conv2D)	(None, 48, 48, 128)	73,856
batch_normalization_4 (BatchNormalization)	(None, 48, 48, 128)	512
conv2d_5 (Conv2D)	(None, 48, 48, 128)	147,584
batch_normalization_5 (BatchNormalization)	(None, 48, 48, 128)	512
max_pooling2d_2 (MaxPooling2D)	(None, 24, 24, 128)	0
dropout_2 (Dropout)	(None, 24, 24, 128)	0
flatten (Flatten)	(None, 73728)	0
dense (Dense)	(None, 256)	18,874,624
batch_normalization_6 (BatchNormalization)	(None, 256)	1,024
dropout_3 (Dropout)	(None, 256)	0
dense_1 (Dense)	(None, 128)	32,896

batch_normalization_7 (BatchNormalization)	(None, 128)	512
dropout_4 (Dropout)	(None, 128)	0
dense_2 (Dense)	(None, 3)	387

Total params: 19,198,243 (73.24 MB)

Trainable params: 19,196,579 (73.23 MB)

Non-trainable params: 1,664 (6.50 KB)

```
In [ ]: # Setup callbacks
def setup_callbacks():
    early_stopping = EarlyStopping(
        monitor='val_accuracy',
        mode='max',
        min_delta=1e-3,
        patience=8,
        restore_best_weights=True,
        verbose=1
    )

    reduce_lr = ReduceLROnPlateau(
        monitor='val_accuracy',
        factor=0.5,
        patience=5,
        min_lr=1e-7,
        verbose=1
    )
    return [early_stopping, reduce_lr]
```

```
In [ ]: def save_model(model, model_name, model_history, test_results):
    import json
    import os
    os.makedirs("out", exist_ok=True)
    os.makedirs("out/model", exist_ok=True)
    os.makedirs("out/model_history", exist_ok=True)
    os.makedirs("out/test_result", exist_ok=True)

    model.save(f"out/model/{model_name}_model.h5")
    history_dict = getattr(model_history, "history", model_history)
    with open(f"out/model_history/{model_name}_history.json", "w") as f:
        json.dump(history_dict, f, indent=2)

    with open(f"out/test_result/{model_name}_test_result.json", "w") as f:
        json.dump(test_results, f, indent=2)
```

```
In [ ]: # Train the baseline model
print("Training Baseline CNN...")
baseline_history = baseline_model.fit(
    X_train_split, y_train_split,
    validation_data=(X_val, y_val),
```

```
    epochs=30,  
    batch_size=32,  
    callbacks=setup_callbacks(),  
    verbose=1  
)
```

Training Baseline CNN...

Epoch 1/30

```
2025-10-30 17:42:44.411230: I external/local_xla/xla/service/service.cc:163] XL  
A service 0x71b7440036f0 initialized for platform CUDA (this does not guarantee  
that XLA will be used). Devices:  
2025-10-30 17:42:44.411265: I external/local_xla/xla/service/service.cc:171]  
StreamExecutor device (0): NVIDIA GeForce RTX 4090, Compute Capability 8.9  
2025-10-30 17:42:44.505560: I tensorflow/compiler/mlir/tensorflow/utils/dump_ml  
ir_util.cc:269] disabling MLIR crash reproducer, set env var `MLIR_CRASH_REPROD  
UCER_DIRECTORY` to enable.  
2025-10-30 17:42:44.886361: I external/local_xla/xla/stream_executor/cuda/cud  
a_dnn.cc:473] Loaded cuDNN version 91301  
2025-10-30 17:42:45.184248: I external/local_xla/xla/service/gpu/autotuning/do  
t_search_space.cc:208] All configs were filtered out because none of them suffi  
ciently match the hints. Maybe the hints set does not contain a good representa  
tive set of valid configs? Working around this by using the full hints set inst  
ead.  
2025-10-30 17:42:45.184589: I external/local_xla/xla/service/gpu/autotuning/do  
t_search_space.cc:208] All configs were filtered out because none of them suffi  
ciently match the hints. Maybe the hints set does not contain a good representa  
tive set of valid configs? Working around this by using the full hints set inst  
ead.  
2025-10-30 17:42:45.184827: I external/local_xla/xla/service/gpu/autotuning/do  
t_search_space.cc:208] All configs were filtered out because none of them suffi  
ciently match the hints. Maybe the hints set does not contain a good representa  
tive set of valid configs? Working around this by using the full hints set inst  
ead.  
2025-10-30 17:42:45.184846: I external/local_xla/xla/service/gpu/autotuning/do  
t_search_space.cc:208] All configs were filtered out because none of them suffi  
ciently match the hints. Maybe the hints set does not contain a good representa  
tive set of valid configs? Working around this by using the full hints set inst  
ead.  
2025-10-30 17:42:45.498916: I external/local_xla/xla/stream_executor/cuda/subpr  
ocess_compilation.cc:346] ptxas warning : Registers are spilled to local memory  
in function 'gemm_fusion_dot_2882', 16 bytes spill stores, 16 bytes spill loads  
  
2025-10-30 17:42:45.929265: I external/local_xla/xla/stream_executor/cuda/subpr  
ocess_compilation.cc:346] ptxas warning : Registers are spilled to local memory  
in function 'gemm_fusion_dot_5058', 520 bytes spill stores, 520 bytes spill loa  
ds  
  
2025-10-30 17:42:45.990900: I external/local_xla/xla/stream_executor/cuda/subpr  
ocess_compilation.cc:346] ptxas warning : Registers are spilled to local memory  
in function 'gemm_fusion_dot_5075', 520 bytes spill stores, 520 bytes spill loa  
ds
```

3/274 ————— 13s 51ms/step - accuracy: 0.3854 - loss: 1.5619 -  
precision: 0.3870 - recall: 0.3438

```
I0000 00:00:1761860571.217336  32196 device_compiler.h:196] Compiled cluster u
sing XLA!  This line is logged at most once for the lifetime of the process.
273/274 0s 21ms/step - accuracy: 0.6530 - loss: 1.0029 - p
recision: 0.6720 - recall: 0.6266

2025-10-30 17:42:57.471757: I external/local_xla/xla/service/gpu/autotuning/do
t_search_space.cc:208] All configs were filtered out because none of them suffi
ciently match the hints. Maybe the hints set does not contain a good representa
tive set of valid configs? Working around this by using the full hints set inst
ead.

2025-10-30 17:42:57.471803: I external/local_xla/xla/service/gpu/autotuning/do
t_search_space.cc:208] All configs were filtered out because none of them suffi
ciently match the hints. Maybe the hints set does not contain a good representa
tive set of valid configs? Working around this by using the full hints set inst
ead.

2025-10-30 17:42:57.659140: I external/local_xla/xla/stream_executor/cuda/subpr
ocess_compilation.cc:346] ptxas warning : Registers are spilled to local memory
in function 'gemm_fusion_dot_4468', 4 bytes spill stores, 4 bytes spill loads

2025-10-30 17:42:57.751806: I external/local_xla/xla/stream_executor/cuda/subpr
ocess_compilation.cc:346] ptxas warning : Registers are spilled to local memory
in function 'gemm_fusion_dot_2882', 16 bytes spill stores, 16 bytes spill loads
```

**274/274** 24s 53ms/step - accuracy: 0.7491 - loss: 0.7348 - precision: 0.7658 - recall: 0.7286 - val\_accuracy: 0.2581 - val\_loss: 6.4320 - val\_precision: 0.2581 - val\_recall: 0.2581 - learning\_rate: 0.0010  
Epoch 2/30

**274/274** 6s 22ms/step - accuracy: 0.8657 - loss: 0.3750 - precision: 0.8719 - recall: 0.8570 - val\_accuracy: 0.3693 - val\_loss: 2.2272 - val\_precision: 0.3681 - val\_recall: 0.3376 - learning\_rate: 0.0010  
Epoch 3/30

**274/274** 10s 22ms/step - accuracy: 0.8848 - loss: 0.3162 - precision: 0.8888 - recall: 0.8771 - val\_accuracy: 0.7154 - val\_loss: 1.2481 - val\_precision: 0.7154 - val\_recall: 0.7154 - learning\_rate: 0.0010  
Epoch 4/30

**274/274** 6s 23ms/step - accuracy: 0.9205 - loss: 0.2203 - precision: 0.9240 - recall: 0.9168 - val\_accuracy: 0.8732 - val\_loss: 0.3465 - val\_precision: 0.8770 - val\_recall: 0.8668 - learning\_rate: 0.0010  
Epoch 5/30

**274/274** 11s 25ms/step - accuracy: 0.9298 - loss: 0.1966 - precision: 0.9322 - recall: 0.9274 - val\_accuracy: 0.7225 - val\_loss: 0.8721 - val\_precision: 0.7267 - val\_recall: 0.7206 - learning\_rate: 0.0010  
Epoch 6/30

**274/274** 11s 23ms/step - accuracy: 0.9404 - loss: 0.1698 - precision: 0.9431 - recall: 0.9393 - val\_accuracy: 0.8260 - val\_loss: 0.7025 - val\_precision: 0.8306 - val\_recall: 0.8182 - learning\_rate: 0.0010  
Epoch 7/30

**274/274** 6s 23ms/step - accuracy: 0.9498 - loss: 0.1425 - precision: 0.9508 - recall: 0.9478 - val\_accuracy: 0.9017 - val\_loss: 0.2606 - val\_precision: 0.9057 - val\_recall: 0.9004 - learning\_rate: 0.0010  
Epoch 8/30

**274/274** 6s 21ms/step - accuracy: 0.9560 - loss: 0.1206 - precision: 0.9575 - recall: 0.9541 - val\_accuracy: 0.6326 - val\_loss: 1.3726 - val\_precision: 0.6324 - val\_recall: 0.6320 - learning\_rate: 0.0010  
Epoch 9/30

**274/274** 7s 24ms/step - accuracy: 0.9551 - loss: 0.1214 - precision: 0.9566 - recall: 0.9541 - val\_accuracy: 0.7768 - val\_loss: 0.7604 - val\_precision: 0.7808 - val\_recall: 0.7743 - learning\_rate: 0.0010  
Epoch 10/30

**274/274** 6s 21ms/step - accuracy: 0.9544 - loss: 0.1240 - precision: 0.9558 - recall: 0.9527 - val\_accuracy: 0.7400 - val\_loss: 1.3113 - val\_precision: 0.7430 - val\_recall: 0.7367 - learning\_rate: 0.0010  
Epoch 11/30

**274/274** 6s 20ms/step - accuracy: 0.9511 - loss: 0.1371 - precision: 0.9523 - recall: 0.9499 - val\_accuracy: 0.8732 - val\_loss: 0.5964 - val\_precision: 0.8787 - val\_recall: 0.8719 - learning\_rate: 0.0010  
Epoch 12/30

**273/274** 0s 20ms/step - accuracy: 0.9685 - loss: 0.0961 - precision: 0.9689 - recall: 0.9675  
Epoch 12: ReduceLROnPlateau reducing learning rate to 0.0005000000237487257.

**274/274** 6s 21ms/step - accuracy: 0.9636 - loss: 0.1028 - precision: 0.9648 - recall: 0.9621 - val\_accuracy: 0.8299 - val\_loss: 0.6632 - val\_precision: 0.8305 - val\_recall: 0.8273 - learning\_rate: 0.0010  
Epoch 13/30

**274/274** 6s 22ms/step - accuracy: 0.9743 - loss: 0.0736 - precision: 0.9744 - recall: 0.9727 - val\_accuracy: 0.9580 - val\_loss: 0.1680 - val\_precision: 0.9592 - val\_recall: 0.9580 - learning\_rate: 5.0000e-04

Epoch 14/30  
**274/274** 6s 21ms/step - accuracy: 0.9793 - loss: 0.0582 - precision: 0.9796 - recall: 0.9792 - val\_accuracy: 0.9515 - val\_loss: 0.1763 - val\_precision: 0.9521 - val\_recall: 0.9508 - learning\_rate: 5.0000e-04  
Epoch 15/30  
**274/274** 7s 24ms/step - accuracy: 0.9801 - loss: 0.0576 - precision: 0.9811 - recall: 0.9796 - val\_accuracy: 0.9444 - val\_loss: 0.2188 - val\_precision: 0.9456 - val\_recall: 0.9437 - learning\_rate: 5.0000e-04  
Epoch 16/30  
**274/274** 6s 21ms/step - accuracy: 0.9818 - loss: 0.0477 - precision: 0.9821 - recall: 0.9816 - val\_accuracy: 0.9528 - val\_loss: 0.1733 - val\_precision: 0.9534 - val\_recall: 0.9521 - learning\_rate: 5.0000e-04  
Epoch 17/30  
**274/274** 6s 22ms/step - accuracy: 0.9752 - loss: 0.0689 - precision: 0.9754 - recall: 0.9745 - val\_accuracy: 0.7956 - val\_loss: 0.8831 - val\_precision: 0.7959 - val\_recall: 0.7943 - learning\_rate: 5.0000e-04  
Epoch 18/30  
**274/274** 0s 20ms/step - accuracy: 0.9817 - loss: 0.0557 - precision: 0.9820 - recall: 0.9811  
Epoch 18: ReduceLROnPlateau reducing learning rate to 0.0002500000118743628.  
**274/274** 6s 21ms/step - accuracy: 0.9777 - loss: 0.0644 - precision: 0.9781 - recall: 0.9772 - val\_accuracy: 0.9211 - val\_loss: 0.2837 - val\_precision: 0.9223 - val\_recall: 0.9211 - learning\_rate: 5.0000e-04  
Epoch 19/30  
**274/274** 6s 21ms/step - accuracy: 0.9812 - loss: 0.0516 - precision: 0.9815 - recall: 0.9804 - val\_accuracy: 0.9644 - val\_loss: 0.1347 - val\_precision: 0.9657 - val\_recall: 0.9638 - learning\_rate: 2.5000e-04  
Epoch 20/30  
**274/274** 7s 25ms/step - accuracy: 0.9881 - loss: 0.0362 - precision: 0.9883 - recall: 0.9876 - val\_accuracy: 0.9644 - val\_loss: 0.1315 - val\_precision: 0.9650 - val\_recall: 0.9638 - learning\_rate: 2.5000e-04  
Epoch 21/30  
**274/274** 6s 21ms/step - accuracy: 0.9904 - loss: 0.0261 - precision: 0.9906 - recall: 0.9902 - val\_accuracy: 0.9560 - val\_loss: 0.1822 - val\_precision: 0.9560 - val\_recall: 0.9554 - learning\_rate: 2.5000e-04  
Epoch 22/30  
**274/274** 6s 21ms/step - accuracy: 0.9928 - loss: 0.0238 - precision: 0.9931 - recall: 0.9925 - val\_accuracy: 0.9644 - val\_loss: 0.1665 - val\_precision: 0.9657 - val\_recall: 0.9638 - learning\_rate: 2.5000e-04  
Epoch 23/30  
**274/274** 6s 21ms/step - accuracy: 0.9929 - loss: 0.0203 - precision: 0.9930 - recall: 0.9929 - val\_accuracy: 0.9036 - val\_loss: 0.4909 - val\_precision: 0.9041 - val\_recall: 0.9030 - learning\_rate: 2.5000e-04  
Epoch 24/30  
**273/274** 0s 21ms/step - accuracy: 0.9928 - loss: 0.0232 - precision: 0.9930 - recall: 0.9928  
Epoch 24: ReduceLROnPlateau reducing learning rate to 0.0001250000059371814.  
**274/274** 6s 22ms/step - accuracy: 0.9928 - loss: 0.0240 - precision: 0.9929 - recall: 0.9928 - val\_accuracy: 0.9599 - val\_loss: 0.1861 - val\_precision: 0.9605 - val\_recall: 0.9592 - learning\_rate: 2.5000e-04  
Epoch 25/30  
**274/274** 6s 21ms/step - accuracy: 0.9937 - loss: 0.0195 - precision: 0.9943 - recall: 0.9935 - val\_accuracy: 0.9670 - val\_loss: 0.1572 - val\_precision: 0.9676 - val\_recall: 0.9670 - learning\_rate: 1.2500e-04

```

Epoch 26/30
274/274 7s 24ms/step - accuracy: 0.9928 - loss: 0.0221 - p
recision: 0.9929 - recall: 0.9928 - val_accuracy: 0.9618 - val_loss: 0.1617 - v
al_precision: 0.9618 - val_recall: 0.9612 - learning_rate: 1.2500e-04
Epoch 27/30
274/274 6s 21ms/step - accuracy: 0.9933 - loss: 0.0193 - p
recision: 0.9933 - recall: 0.9930 - val_accuracy: 0.9618 - val_loss: 0.1636 - v
al_precision: 0.9637 - val_recall: 0.9618 - learning_rate: 1.2500e-04
Epoch 28/30
274/274 6s 21ms/step - accuracy: 0.9947 - loss: 0.0169 - p
recision: 0.9947 - recall: 0.9947 - val_accuracy: 0.9696 - val_loss: 0.1473 - v
al_precision: 0.9702 - val_recall: 0.9696 - learning_rate: 1.2500e-04
Epoch 29/30
274/274 10s 21ms/step - accuracy: 0.9961 - loss: 0.0123 -
precision: 0.9961 - recall: 0.9960 - val_accuracy: 0.9696 - val_loss: 0.1319 -
val_precision: 0.9696 - val_recall: 0.9683 - learning_rate: 1.2500e-04
Epoch 30/30
274/274 6s 21ms/step - accuracy: 0.9969 - loss: 0.0110 - p
recision: 0.9969 - recall: 0.9968 - val_accuracy: 0.9722 - val_loss: 0.1515 - v
al_precision: 0.9728 - val_recall: 0.9709 - learning_rate: 1.2500e-04
Restoring model weights from the end of the best epoch: 30.

```

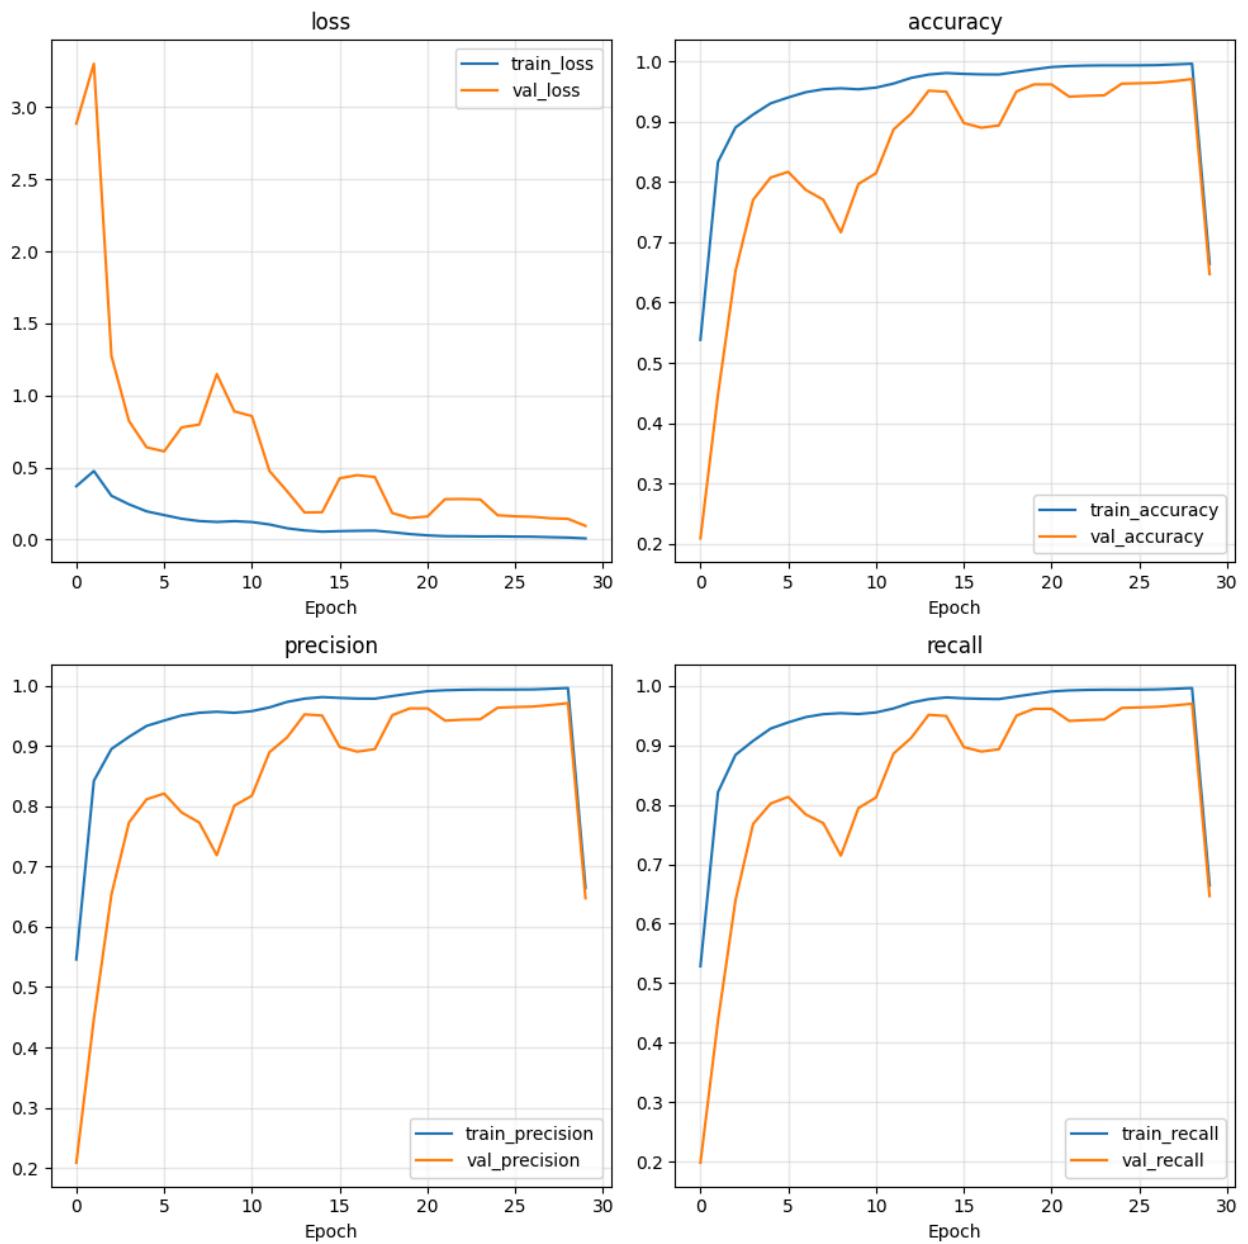
```

In [ ]: # Define a helper function to smooth training curves using a simple moving average
def smooth(xs, w=3):
    if w <= 1: return xs
    xs = np.array(xs, dtype=float)
    return np.convolve(xs, np.ones(w)/w, mode='same')

metrics = [
    ('loss', 'val_loss'),
    ('accuracy', 'val_accuracy'),
    ('precision', 'val_precision'),
    ('recall', 'val_recall'),
]

plt.figure(figsize=(10, 10))
for i, (m, vm) in enumerate(metrics, 1):
    plt.subplot(2, 2, i)
    train_curve = baseline_history.history .get(m, [])
    val_curve   = baseline_history.history .get(vm, [])
    w = 3
    plt.plot(smooth(train_curve, w=w), label=f'train_{m}')
    plt.plot(smooth(val_curve,   w=w), label=f'val_{m}')
    plt.title(m)
    plt.xlabel('Epoch')
    plt.grid(True, alpha=0.3)
    plt.legend()
plt.tight_layout()
plt.show()

```



```
In [ ]: # Evaluate on test set
print("\nBaseline CNN - Test Set Evaluation:")
baseline_test_loss, baseline_test_acc, baseline_test_prec, baseline_test_recal
print(f"Test Accuracy: {baseline_test_acc:.4f}")
print(f"Test Precision: {baseline_test_prec:.4f}")
print(f"Test Recall: {baseline_test_recall:.4f}")
print(f"Test F1-Score: {2 * (baseline_test_prec * baseline_test_recall) / (ba

test_results = {
    "model_name": "Baseline_CNN",
    "test_loss": float(baseline_test_loss),
    "test_accuracy": float(baseline_test_acc),
    "test_precision": float(baseline_test_prec),
    "test_recall": float(baseline_test_recall),
    "test_f1": float(2 * (baseline_test_prec * baseline_test_recall) / (ba
}
```

```
Baseline CNN - Test Set Evaluation:  
150/152 ━━━━━━ 0s 6ms/step - accuracy: 0.9800 - loss: 0.0639 - pr  
ecision: 0.9801 - recall: 0.9796  
2025-10-30 17:46:31.106193: I external/local_xla/xla/service/gpu/autotuning/do  
t_search_space.cc:208] All configs were filtered out because none of them suffi  
ciently match the hints. Maybe the hints set does not contain a good representa  
tive set of valid configs? Working around this by using the full hints set inst  
ead.  
2025-10-30 17:46:31.303204: I external/local_xla/xla/stream_executor/cuda/subpr  
ocess_compilation.cc:346] ptxas warning : Registers are spilled to local memory  
in function 'gemm_fusion_dot_260', 16 bytes spill stores, 16 bytes spill loads  
  
152/152 ━━━━━━ 2s 13ms/step - accuracy: 0.9810 - loss: 0.0715 - p  
recision: 0.9812 - recall: 0.9806  
Test Accuracy: 0.9810  
Test Precision: 0.9812  
Test Recall: 0.9806  
Test F1-Score: 0.9809
```

```
In [ ]: save_model(baseline_model, "baseline_model", baseline_history, test_results)
```

```
WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `k  
eras.saving.save_model(model)`. This file format is considered legacy. We recom  
mend using instead the native Keras format, e.g. `model.save('my_model.keras')`  
or `keras.saving.save_model(model, 'my_model.keras')`.
```

### 3. Transfer Learning with ResNet

Approach:

1. Load ResNet50 without top layers
2. Freeze base layers initially
3. Add custom classification head
4. Fine-tune top layers after initial training

**Fine-tuning:** Unfreeze last 20 layers

```
In [ ]: def create_resnet_model(input_shape=(192, 192, 3), num_classes=3, fine_tune=False  
        ):  
    """  
    Create ResNet50 transfer learning model.  
    """  
    # Load pre-trained ResNet50  
    base_model = ResNet50(  
        weights='imagenet',  
        include_top=False,  
        input_shape=input_shape  
    )  
  
    # Freeze base model layers  
    base_model.trainable = fine_tune
```

```

if fine_tune:
    # Fine-tune from layer 140 onwards
    for layer in base_model.layers[:140]:
        layer.trainable = False

# Build model
model = Sequential([
    base_model,
    GlobalAveragePooling2D(),
    Dense(512, activation='relu'),
    BatchNormalization(),
    Dropout(0.5),
    Dense(256, activation='relu'),
    BatchNormalization(),
    Dropout(0.5),
    Dense(num_classes, activation='softmax')
], name='ResNet50_Transfer')

return model

```

```

In [ ]: # Create ResNet model
resnet_model = create_resnet_model(fine_tune=False)

resnet_model.compile(
    optimizer=Adam(learning_rate=0.0001),
    loss='categorical_crossentropy',
    metrics=['accuracy',
              tf.keras.metrics.Precision(name='precision'),
              tf.keras.metrics.Recall(name='recall')])
)

print(f"Total layers: {len(resnet_model.layers)}")
print(f"Trainable layers: {sum([layer.trainable for layer in resnet_model.layers])}")

```

Total layers: 9  
Trainable layers: 8

```

In [ ]: # Train ResNet model
print("Training ResNet50 (frozen base)...")
resnet_history = resnet_model.fit(
    X_train_split, y_train_split,
    validation_data=(X_val, y_val),
    epochs=30,
    batch_size=32,
    callbacks=setup_callbacks(),
    verbose=1
)

```

Training ResNet50 (frozen base)...  
Epoch 1/30

```
2025-10-30 17:47:28.201459: I external/local_xla/xla/service/gpu/autotuning/do
t_search_space.cc:208] All configs were filtered out because none of them suffi
ciently match the hints. Maybe the hints set does not contain a good representa
tive set of valid configs? Working around this by using the full hints set inst
ead.
2025-10-30 17:47:28.201513: I external/local_xla/xla/service/gpu/autotuning/do
t_search_space.cc:208] All configs were filtered out because none of them suffi
ciently match the hints. Maybe the hints set does not contain a good representa
tive set of valid configs? Working around this by using the full hints set inst
ead.
2025-10-30 17:47:28.201528: I external/local_xla/xla/service/gpu/autotuning/do
t_search_space.cc:208] All configs were filtered out because none of them suffi
ciently match the hints. Maybe the hints set does not contain a good representa
tive set of valid configs? Working around this by using the full hints set inst
ead.
2025-10-30 17:47:28.201534: I external/local_xla/xla/service/gpu/autotuning/do
t_search_space.cc:208] All configs were filtered out because none of them suffi
ciently match the hints. Maybe the hints set does not contain a good representa
tive set of valid configs? Working around this by using the full hints set inst
ead.
2025-10-30 17:47:28.611113: I external/local_xla/xla/stream_executor/cuda/subpr
ocess_compilation.cc:346] ptxas warning : Registers are spilled to local memory
in function 'gemm_fusion_dot_5767', 16 bytes spill stores, 16 bytes spill loads
2025-10-30 17:47:28.745110: I external/local_xla/xla/stream_executor/cuda/subpr
ocess_compilation.cc:346] ptxas warning : Registers are spilled to local memory
in function 'gemm_fusion_dot_6328', 12 bytes spill stores, 12 bytes spill loads
2025-10-30 17:47:28.829541: I external/local_xla/xla/stream_executor/cuda/subpr
ocess_compilation.cc:346] ptxas warning : Registers are spilled to local memory
in function 'gemm_fusion_dot_6514', 4 bytes spill stores, 4 bytes spill loads
2025-10-30 17:47:28.983619: I external/local_xla/xla/stream_executor/cuda/subpr
ocess_compilation.cc:346] ptxas warning : Registers are spilled to local memory
in function 'gemm_fusion_dot_6514', 520 bytes spill stores, 520 bytes spill lo
ds
2025-10-30 17:47:29.008953: I external/local_xla/xla/stream_executor/cuda/subpr
ocess_compilation.cc:346] ptxas warning : Registers are spilled to local memory
in function 'gemm_fusion_dot_6347', 520 bytes spill stores, 520 bytes spill lo
ds
```

**271/274** — **0s** 11ms/step - accuracy: 0.4488 - loss: 1.5123 - p
recision: 0.4626 - recall: 0.4033

2025-10-30 17:47:35.017777: I external/local\_xla/xla/service/gpu/autotuning/do\_t\_search\_space.cc:208] All configs were filtered out because none of them sufficiently match the hints. Maybe the hints set does not contain a good representative set of valid configs? Working around this by using the full hints set instead.

2025-10-30 17:47:35.017823: I external/local\_xla/xla/service/gpu/autotuning/do\_t\_search\_space.cc:208] All configs were filtered out because none of them sufficiently match the hints. Maybe the hints set does not contain a good representative set of valid configs? Working around this by using the full hints set instead.

2025-10-30 17:47:35.303363: I external/local\_xla/xla/stream\_executor/cuda/subprocess\_compilation.cc:346] ptxas warning : Registers are spilled to local memory in function 'gemm\_fusion\_dot\_5767', 16 bytes spill stores, 16 bytes spill loads

2025-10-30 17:47:35.333835: I external/local\_xla/xla/stream\_executor/cuda/subprocess\_compilation.cc:346] ptxas warning : Registers are spilled to local memory in function 'gemm\_fusion\_dot\_6328', 12 bytes spill stores, 12 bytes spill loads

2025-10-30 17:47:35.343869: I external/local\_xla/xla/stream\_executor/cuda/subprocess\_compilation.cc:346] ptxas warning : Registers are spilled to local memory in function 'gemm\_fusion\_dot\_5874', 4 bytes spill stores, 4 bytes spill loads

2025-10-30 17:47:35.461177: I external/local\_xla/xla/stream\_executor/cuda/subprocess\_compilation.cc:346] ptxas warning : Registers are spilled to local memory in function 'gemm\_fusion\_dot\_5767', 8 bytes spill stores, 8 bytes spill loads

**274/274** 18s 38ms/step - accuracy: 0.4953 - loss: 1.3647 - precision: 0.5118 - recall: 0.4498 - val\_accuracy: 0.2497 - val\_loss: 1.1945 - val\_precision: 0.4770 - val\_recall: 0.0873 - learning\_rate: 1.0000e-04  
Epoch 2/30

**274/274** 4s 14ms/step - accuracy: 0.6068 - loss: 1.0289 - precision: 0.6326 - recall: 0.5669 - val\_accuracy: 0.4864 - val\_loss: 0.9195 - val\_precision: 0.4865 - val\_recall: 0.4437 - learning\_rate: 1.0000e-04  
Epoch 3/30

**274/274** 4s 14ms/step - accuracy: 0.6728 - loss: 0.8489 - precision: 0.6969 - recall: 0.6368 - val\_accuracy: 0.1429 - val\_loss: 2.0124 - val\_precision: 0.1416 - val\_recall: 0.1404 - learning\_rate: 1.0000e-04  
Epoch 4/30

**274/274** 4s 14ms/step - accuracy: 0.7191 - loss: 0.7273 - precision: 0.7433 - recall: 0.6867 - val\_accuracy: 0.2691 - val\_loss: 3.3213 - val\_precision: 0.2681 - val\_recall: 0.2665 - learning\_rate: 1.0000e-04  
Epoch 5/30

**274/274** 4s 14ms/step - accuracy: 0.7520 - loss: 0.6581 - precision: 0.7740 - recall: 0.7265 - val\_accuracy: 0.2607 - val\_loss: 5.3386 - val\_precision: 0.2599 - val\_recall: 0.2581 - learning\_rate: 1.0000e-04  
Epoch 6/30

**274/274** 5s 17ms/step - accuracy: 0.7810 - loss: 0.6033 - precision: 0.7966 - recall: 0.7580 - val\_accuracy: 0.5414 - val\_loss: 1.1390 - val\_precision: 0.5483 - val\_recall: 0.5291 - learning\_rate: 1.0000e-04  
Epoch 7/30

**274/274** 5s 13ms/step - accuracy: 0.7900 - loss: 0.5688 - precision: 0.8041 - recall: 0.7687 - val\_accuracy: 0.7070 - val\_loss: 0.9892 - val\_precision: 0.7097 - val\_recall: 0.7070 - learning\_rate: 1.0000e-04  
Epoch 8/30

**274/274** 3s 12ms/step - accuracy: 0.7994 - loss: 0.5354 - precision: 0.8112 - recall: 0.7825 - val\_accuracy: 0.6546 - val\_loss: 0.9649 - val\_precision: 0.6589 - val\_recall: 0.6397 - learning\_rate: 1.0000e-04  
Epoch 9/30

**274/274** 4s 13ms/step - accuracy: 0.8053 - loss: 0.5176 - precision: 0.8179 - recall: 0.7891 - val\_accuracy: 0.2400 - val\_loss: 7.1607 - val\_precision: 0.2400 - val\_recall: 0.2400 - learning\_rate: 1.0000e-04  
Epoch 10/30

**274/274** 4s 14ms/step - accuracy: 0.8187 - loss: 0.4825 - precision: 0.8286 - recall: 0.8037 - val\_accuracy: 0.8260 - val\_loss: 0.4217 - val\_precision: 0.8320 - val\_recall: 0.8169 - learning\_rate: 1.0000e-04  
Epoch 11/30

**274/274** 4s 13ms/step - accuracy: 0.8325 - loss: 0.4518 - precision: 0.8398 - recall: 0.8174 - val\_accuracy: 0.6785 - val\_loss: 1.8122 - val\_precision: 0.6785 - val\_recall: 0.6785 - learning\_rate: 1.0000e-04  
Epoch 12/30

**274/274** 4s 13ms/step - accuracy: 0.8280 - loss: 0.4486 - precision: 0.8384 - recall: 0.8164 - val\_accuracy: 0.6902 - val\_loss: 1.5984 - val\_precision: 0.6915 - val\_recall: 0.6902 - learning\_rate: 1.0000e-04  
Epoch 13/30

**274/274** 4s 14ms/step - accuracy: 0.8322 - loss: 0.4407 - precision: 0.8413 - recall: 0.8223 - val\_accuracy: 0.3137 - val\_loss: 6.4448 - val\_precision: 0.3145 - val\_recall: 0.3137 - learning\_rate: 1.0000e-04  
Epoch 14/30

**274/274** 5s 16ms/step - accuracy: 0.8423 - loss: 0.4127 - precision: 0.8489 - recall: 0.8322 - val\_accuracy: 0.7115 - val\_loss: 0.6412 - val\_precision: 0.7115 - val\_recall: 0.7115 - learning\_rate: 1.0000e-04

```
al_precision: 0.7159 - val_recall: 0.7089 - learning_rate: 1.0000e-04
Epoch 15/30
274/274 0s 11ms/step - accuracy: 0.8419 - loss: 0.4057 - p
recision: 0.8519 - recall: 0.8349
Epoch 15: ReduceLROnPlateau reducing learning rate to 4.999999873689376e-05.
274/274 4s 14ms/step - accuracy: 0.8425 - loss: 0.4065 - p
recision: 0.8512 - recall: 0.8354 - val_accuracy: 0.7840 - val_loss: 0.7319 - v
al_precision: 0.7847 - val_recall: 0.7827 - learning_rate: 1.0000e-04
Epoch 16/30
274/274 4s 14ms/step - accuracy: 0.8492 - loss: 0.3874 - p
recision: 0.8578 - recall: 0.8405 - val_accuracy: 0.8441 - val_loss: 0.4366 - v
al_precision: 0.8570 - val_recall: 0.8221 - learning_rate: 5.0000e-05
Epoch 17/30
274/274 4s 14ms/step - accuracy: 0.8489 - loss: 0.3809 - p
recision: 0.8569 - recall: 0.8398 - val_accuracy: 0.8687 - val_loss: 0.3502 - v
al_precision: 0.8737 - val_recall: 0.8635 - learning_rate: 5.0000e-05
Epoch 18/30
274/274 5s 14ms/step - accuracy: 0.8509 - loss: 0.3812 - p
recision: 0.8582 - recall: 0.8416 - val_accuracy: 0.7296 - val_loss: 0.7859 - v
al_precision: 0.7303 - val_recall: 0.7270 - learning_rate: 5.0000e-05
Epoch 19/30
274/274 4s 13ms/step - accuracy: 0.8565 - loss: 0.3752 - p
recision: 0.8622 - recall: 0.8475 - val_accuracy: 0.7432 - val_loss: 0.6245 - v
al_precision: 0.7495 - val_recall: 0.7335 - learning_rate: 5.0000e-05
Epoch 20/30
274/274 4s 14ms/step - accuracy: 0.8606 - loss: 0.3605 - p
recision: 0.8663 - recall: 0.8532 - val_accuracy: 0.8616 - val_loss: 0.3903 - v
al_precision: 0.8636 - val_recall: 0.8603 - learning_rate: 5.0000e-05
Epoch 21/30
274/274 4s 14ms/step - accuracy: 0.8567 - loss: 0.3545 - p
recision: 0.8650 - recall: 0.8486 - val_accuracy: 0.6158 - val_loss: 0.9737 - v
al_precision: 0.6167 - val_recall: 0.6151 - learning_rate: 5.0000e-05
Epoch 22/30
274/274 0s 13ms/step - accuracy: 0.8622 - loss: 0.3565 - p
recision: 0.8667 - recall: 0.8499
Epoch 22: ReduceLROnPlateau reducing learning rate to 2.499999936844688e-05.
274/274 6s 16ms/step - accuracy: 0.8607 - loss: 0.3582 - p
recision: 0.8656 - recall: 0.8520 - val_accuracy: 0.3907 - val_loss: 1.9761 - v
al_precision: 0.3915 - val_recall: 0.3829 - learning_rate: 5.0000e-05
Epoch 23/30
274/274 5s 14ms/step - accuracy: 0.8673 - loss: 0.3442 - p
recision: 0.8735 - recall: 0.8597 - val_accuracy: 0.8765 - val_loss: 0.3135 - v
al_precision: 0.8802 - val_recall: 0.8700 - learning_rate: 2.5000e-05
Epoch 24/30
274/274 4s 14ms/step - accuracy: 0.8660 - loss: 0.3401 - p
recision: 0.8718 - recall: 0.8585 - val_accuracy: 0.7367 - val_loss: 0.7656 - v
al_precision: 0.7365 - val_recall: 0.7342 - learning_rate: 2.5000e-05
Epoch 25/30
274/274 4s 14ms/step - accuracy: 0.8615 - loss: 0.3454 - p
recision: 0.8671 - recall: 0.8549 - val_accuracy: 0.8771 - val_loss: 0.3038 - v
al_precision: 0.8843 - val_recall: 0.8752 - learning_rate: 2.5000e-05
Epoch 26/30
274/274 4s 13ms/step - accuracy: 0.8608 - loss: 0.3456 - p
recision: 0.8656 - recall: 0.8533 - val_accuracy: 0.8344 - val_loss: 0.3995 - v
```

```
al_precision: 0.8392 - val_recall: 0.8305 - learning_rate: 2.5000e-05
Epoch 27/30
274/274 ————— 4s 13ms/step - accuracy: 0.8681 - loss: 0.3352 - p
recision: 0.8752 - recall: 0.8606 - val_accuracy: 0.7865 - val_loss: 0.5376 - v
al_precision: 0.7958 - val_recall: 0.7814 - learning_rate: 2.5000e-05
Epoch 28/30
274/274 ————— 4s 14ms/step - accuracy: 0.8719 - loss: 0.3271 - p
recision: 0.8783 - recall: 0.8657 - val_accuracy: 0.8984 - val_loss: 0.2833 - v
al_precision: 0.9022 - val_recall: 0.8946 - learning_rate: 2.5000e-05
Epoch 29/30
274/274 ————— 4s 13ms/step - accuracy: 0.8678 - loss: 0.3310 - p
recision: 0.8736 - recall: 0.8616 - val_accuracy: 0.8448 - val_loss: 0.4038 - v
al_precision: 0.8621 - val_recall: 0.8292 - learning_rate: 2.5000e-05
Epoch 30/30
274/274 ————— 4s 16ms/step - accuracy: 0.8702 - loss: 0.3317 - p
recision: 0.8756 - recall: 0.8638 - val_accuracy: 0.8952 - val_loss: 0.2821 - v
al_precision: 0.9061 - val_recall: 0.8862 - learning_rate: 2.5000e-05
Restoring model weights from the end of the best epoch: 28.
```

```
In [ ]: # Fine-tune the model
print("\nFine-tuning ResNet50...")

# Unfreeze last layers
for layer in resnet_model.layers[0].layers[-20:]:
    layer.trainable = True

# Recompile with lower learning rate
resnet_model.compile(
    optimizer=Adam(learning_rate=0.00001),
    loss='categorical_crossentropy',
    metrics=['accuracy',
              tf.keras.metrics.Precision(name='precision'),
              tf.keras.metrics.Recall(name='recall')])
)

resnet_history_finetune = resnet_model.fit(
    X_train_split, y_train_split,
    validation_data=(X_val, y_val),
    epochs=15,
    batch_size=32,
    callbacks=setup_callbacks(),
    verbose=1
)
```

Fine-tuning ResNet50...

Epoch 1/15

**274/274** 19s 41ms/step - accuracy: 0.8139 - loss: 0.4967 - precision: 0.8225 - recall: 0.8049 - val\_accuracy: 0.1514 - val\_loss: 1.9722 - val\_precision: 0.1487 - val\_recall: 0.1442 - learning\_rate: 1.0000e-05

Epoch 2/15

**274/274** 5s 18ms/step - accuracy: 0.8647 - loss: 0.3569 - precision: 0.8711 - recall: 0.8566 - val\_accuracy: 0.8169 - val\_loss: 0.4865 - val\_precision: 0.8374 - val\_recall: 0.8027 - learning\_rate: 1.0000e-05

Epoch 3/15

**274/274** 4s 15ms/step - accuracy: 0.8940 - loss: 0.2923 - precision: 0.8986 - recall: 0.8870 - val\_accuracy: 0.8816 - val\_loss: 0.3171 - val\_precision: 0.8838 - val\_recall: 0.8758 - learning\_rate: 1.0000e-05

Epoch 4/15

**274/274** 4s 16ms/step - accuracy: 0.8971 - loss: 0.2663 - precision: 0.9032 - recall: 0.8929 - val\_accuracy: 0.9185 - val\_loss: 0.2386 - val\_precision: 0.9207 - val\_recall: 0.9166 - learning\_rate: 1.0000e-05

Epoch 5/15

**274/274** 5s 15ms/step - accuracy: 0.8999 - loss: 0.2670 - precision: 0.9056 - recall: 0.8958 - val\_accuracy: 0.9075 - val\_loss: 0.2367 - val\_precision: 0.9160 - val\_recall: 0.9030 - learning\_rate: 1.0000e-05

Epoch 6/15

**274/274** 4s 15ms/step - accuracy: 0.9109 - loss: 0.2373 - precision: 0.9152 - recall: 0.9069 - val\_accuracy: 0.8635 - val\_loss: 0.3646 - val\_precision: 0.8654 - val\_recall: 0.8609 - learning\_rate: 1.0000e-05

Epoch 7/15

**274/274** 4s 15ms/step - accuracy: 0.9202 - loss: 0.2159 - precision: 0.9235 - recall: 0.9148 - val\_accuracy: 0.9088 - val\_loss: 0.2567 - val\_precision: 0.9215 - val\_recall: 0.9036 - learning\_rate: 1.0000e-05

Epoch 8/15

**274/274** 4s 16ms/step - accuracy: 0.9257 - loss: 0.2009 - precision: 0.9283 - recall: 0.9211 - val\_accuracy: 0.8596 - val\_loss: 0.4092 - val\_precision: 0.8612 - val\_recall: 0.8551 - learning\_rate: 1.0000e-05

Epoch 9/15

**271/274** 0s 16ms/step - accuracy: 0.9293 - loss: 0.1940 - precision: 0.9336 - recall: 0.9258

Epoch 9: ReduceLROnPlateau reducing learning rate to 4.999999873689376e-06.

**274/274** 5s 18ms/step - accuracy: 0.9292 - loss: 0.1975 - precision: 0.9324 - recall: 0.9264 - val\_accuracy: 0.9185 - val\_loss: 0.2233 - val\_precision: 0.9231 - val\_recall: 0.9166 - learning\_rate: 1.0000e-05

Epoch 10/15

**274/274** 4s 16ms/step - accuracy: 0.9437 - loss: 0.1566 - precision: 0.9467 - recall: 0.9413 - val\_accuracy: 0.9314 - val\_loss: 0.1894 - val\_precision: 0.9326 - val\_recall: 0.9301 - learning\_rate: 5.0000e-06

Epoch 11/15

**274/274** 4s 16ms/step - accuracy: 0.9454 - loss: 0.1529 - precision: 0.9472 - recall: 0.9426 - val\_accuracy: 0.9392 - val\_loss: 0.1659 - val\_precision: 0.9403 - val\_recall: 0.9379 - learning\_rate: 5.0000e-06

Epoch 12/15

**274/274** 4s 16ms/step - accuracy: 0.9535 - loss: 0.1322 - precision: 0.9559 - recall: 0.9515 - val\_accuracy: 0.9334 - val\_loss: 0.1781 - val\_precision: 0.9351 - val\_recall: 0.9314 - learning\_rate: 5.0000e-06

Epoch 13/15

**274/274** 4s 15ms/step - accuracy: 0.9535 - loss: 0.1318 - p

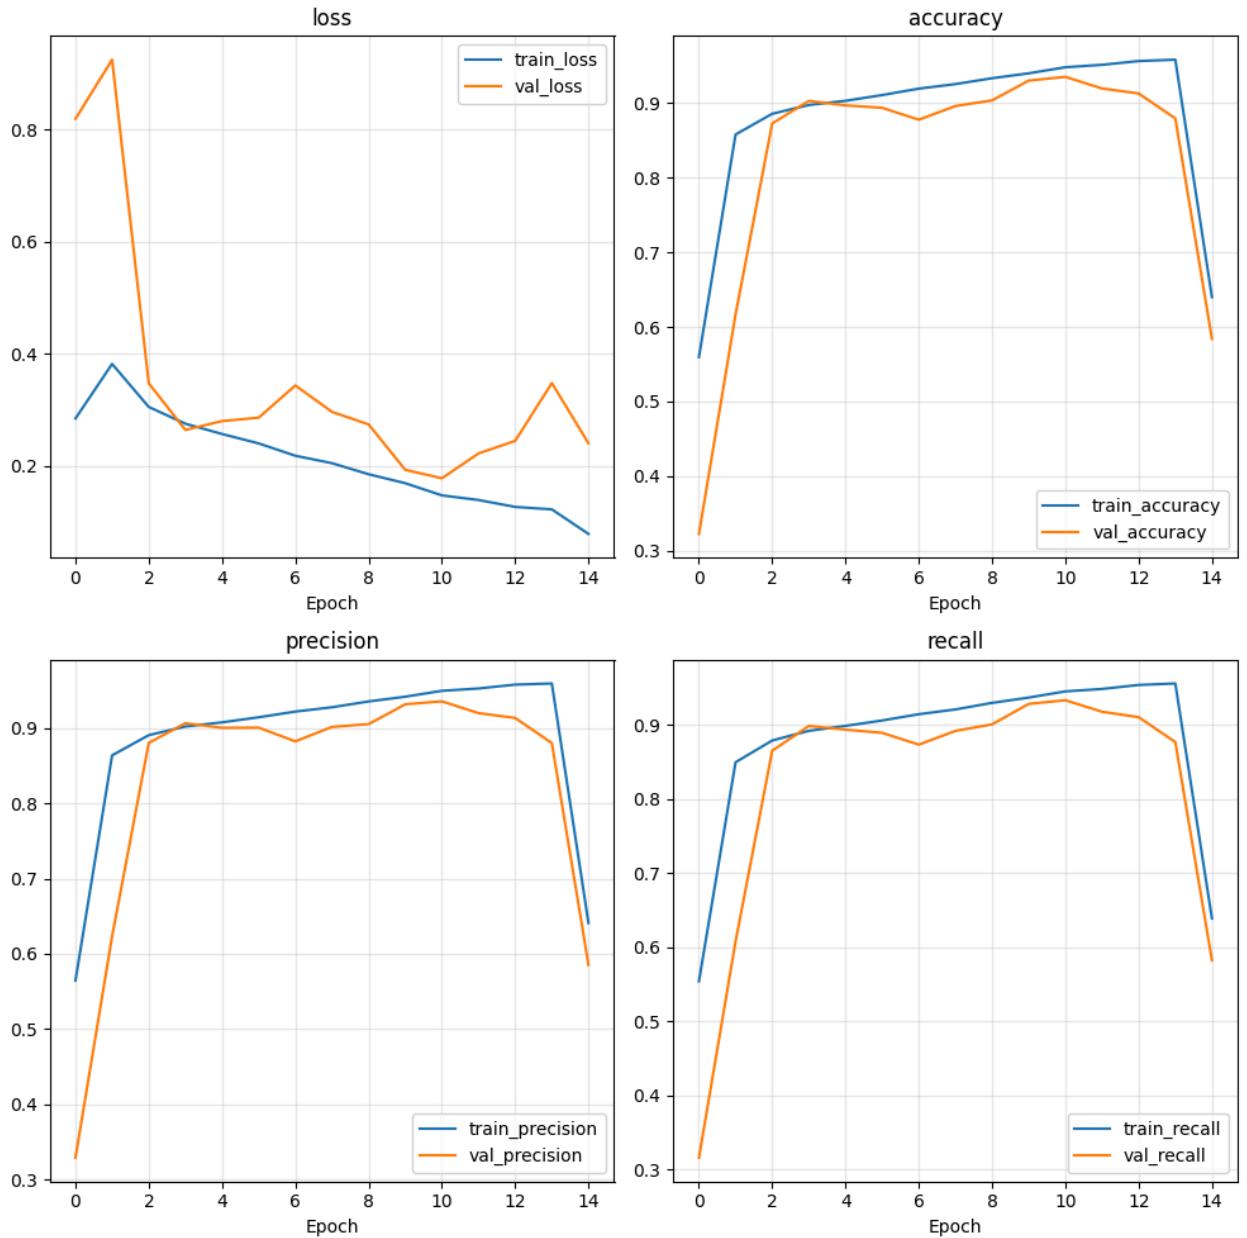
```
precision: 0.9560 - recall: 0.9511 - val_accuracy: 0.8849 - val_loss: 0.3226 - val_precision: 0.8853 - val_recall: 0.8836 - learning_rate: 5.0000e-06
Epoch 14/15
274/274 ————— 4s 15ms/step - accuracy: 0.9605 - loss: 0.1159 - precision: 0.9626 - recall: 0.9588 - val_accuracy: 0.9191 - val_loss: 0.2331 - val_precision: 0.9213 - val_recall: 0.9159 - learning_rate: 5.0000e-06
Epoch 15/15
274/274 ————— 4s 16ms/step - accuracy: 0.9591 - loss: 0.1191 - precision: 0.9604 - recall: 0.9573 - val_accuracy: 0.8331 - val_loss: 0.4879 - val_precision: 0.8349 - val_recall: 0.8305 - learning_rate: 5.0000e-06
Restoring model weights from the end of the best epoch: 11.
```

In [ ]: *#helper function used previously*

```
def smooth(xs, w=3):
    if w <= 1: return xs
    xs = np.array(xs, dtype=float)
    return np.convolve(xs, np.ones(w)/w, mode='same')

metrics = [
    ('loss', 'val_loss'),
    ('accuracy', 'val_accuracy'),
    ('precision', 'val_precision'),
    ('recall', 'val_recall'),
]

plt.figure(figsize=(10, 10))
for i, (m, vm) in enumerate(metrics, 1):
    plt.subplot(2, 2, i)
    train_curve = resnet_history_finetune.history .get(m, [])
    val_curve   = resnet_history_finetune.history .get(vm, [])
    w = 3
    plt.plot(smooth(train_curve, w=w), label=f'train_{m}')
    plt.plot(smooth(val_curve,   w=w), label=f'val_{m}')
    plt.title(m)
    plt.xlabel('Epoch')
    plt.grid(True, alpha=0.3)
    plt.legend()
plt.tight_layout()
plt.show()
```



```
In [ ]: # Evaluate ResNet on test set
print("\nResNet50 - Test Set Evaluation:")
resnet_test_loss, resnet_test_acc, resnet_test_prec, resnet_test_recall = resnet.evaluate()
print(f"Test Accuracy: {resnet_test_acc:.4f}")
print(f"Test Precision: {resnet_test_prec:.4f}")
print(f"Test Recall: {resnet_test_recall:.4f}")
print(f"Test F1-Score: {2 * (resnet_test_prec * resnet_test_recall) / (resnet_test_prec + resnet_test_recall)}")

test_results = {
    "model_name": "ResNet50",
    "test_loss": float(resnet_test_loss),
    "test_accuracy": float(resnet_test_acc),
    "test_precision": float(resnet_test_prec),
    "test_recall": float(resnet_test_recall),
    "test_f1": float(2 * (resnet_test_prec * resnet_test_recall) / (resnet_test_prec + resnet_test_recall))
}
```

```

ResNet50 - Test Set Evaluation:
151/152 ━━━━━━ 0s 11ms/step - accuracy: 0.9418 - loss: 0.1593 - p
recision: 0.9428 - recall: 0.9404
2025-10-30 17:51:23.926779: I external/local_xla/xla/service/gpu/autotuning/do
t_search_space.cc:208] All configs were filtered out because none of them suffi
ciently match the hints. Maybe the hints set does not contain a good representa
tive set of valid configs? Working around this by using the full hints set inst
ead.
2025-10-30 17:51:24.080381: I external/local_xla/xla/stream_executor/cuda/subpr
ocess_compilation.cc:346] ptxas warning : Registers are spilled to local memory
in function 'gemm_fusion_dot_1798', 4 bytes spill stores, 4 bytes spill loads

2025-10-30 17:51:24.158708: I external/local_xla/xla/stream_executor/cuda/subpr
ocess_compilation.cc:346] ptxas warning : Registers are spilled to local memory
in function 'gemm_fusion_dot_1787', 16 bytes spill stores, 16 bytes spill loads

2025-10-30 17:51:24.297746: I external/local_xla/xla/stream_executor/cuda/subpr
ocess_compilation.cc:346] ptxas warning : Registers are spilled to local memory
in function 'gemm_fusion_dot_1787', 8 bytes spill stores, 8 bytes spill loads

152/152 ━━━━━━ 4s 23ms/step - accuracy: 0.9338 - loss: 0.1850 - p
recision: 0.9357 - recall: 0.9326
Test Accuracy: 0.9338
Test Precision: 0.9357
Test Recall: 0.9326
Test F1-Score: 0.9341

```

```
In [ ]: save_model(resnet_model, "ResNet50", resnet_history_finetune, test_results)
```

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `k
eras.saving.save\_model(model)`. This file format is considered legacy. We recom
mend using instead the native Keras format, e.g. `model.save('my\_model.keras')` 
or `keras.saving.save\_model(model, 'my\_model.keras')`.

## Impact of Pre-trained Features

Benefits:

1. Faster Convergence: Pre-trained features accelerate training
2. Better Generalization: ImageNet features transfer well to medical
images
3. Lower Data Requirements: Effective even with limited training data
4. Improved Performance: Typically outperforms models trained from
scratch

## 4. Additional Architectures

Implement three additional models:

1. **VGG19** - Deep architecture with small filters

2. **InceptionV3** - Multi-scale feature extraction
3. **mobilenet** - Faster and light

## Model 1: VGG19

```
In [ ]: def create_vgg19_model(input_shape=(192, 192, 3), num_classes=3):
    """
    Create VGG19 transfer learning model.
    """
    base_model = VGG19(
        weights='imagenet',
        include_top=False,
        input_shape=input_shape
    )

    base_model.trainable = False

    model = Sequential([
        base_model,
        GlobalAveragePooling2D(),
        Dense(512, activation='relu'),
        BatchNormalization(),
        Dropout(0.5),
        Dense(256, activation='relu'),
        BatchNormalization(),
        Dropout(0.5),
        Dense(num_classes, activation='softmax')
    ], name='VGG19_Transfer')

    return model

# Create and compile VGG19
vgg19_model = create_vgg19_model()
vgg19_model.compile(
    optimizer=Adam(learning_rate=0.0001),
    loss='categorical_crossentropy',
    metrics=['accuracy',
              tf.keras.metrics.Precision(name='precision'),
              tf.keras.metrics.Recall(name='recall')])
)

print("VGG19 Model Created")
```

VGG19 Model Created

```
In [ ]: # Train VGG19
print("Training VGG19...")
vgg19_history = vgg19_model.fit(
    X_train_split, y_train_split,
    validation_data=(X_val, y_val),
    epochs=25,
    batch_size=32,
```

```
    callbacks=setup_callbacks(),
    verbose=1
)
```

Training VGG19...

Epoch 1/25

```
2025-10-30 17:51:49.683271: I external/local_xla/xla/service/gpu/autotuning/do
t_search_space.cc:208] All configs were filtered out because none of them suffi
ciently match the hints. Maybe the hints set does not contain a good representa
tive set of valid configs? Working around this by using the full hints set inst
ead.
```

```
2025-10-30 17:51:49.889475: I external/local_xla/xla/stream_executor/cuda/subpr
ocess_compilation.cc:346] ptxas warning : Registers are spilled to local memory
in function 'gemm_fusion_dot_2438', 4 bytes spill stores, 4 bytes spill loads
```

```
2025-10-30 17:51:50.029843: I external/local_xla/xla/stream_executor/cuda/subpr
ocess_compilation.cc:346] ptxas warning : Registers are spilled to local memory
in function 'gemm_fusion_dot_2438', 520 bytes spill stores, 520 bytes spill loa
ds
```

**274/274** 16s 40ms/step - accuracy: 0.5963 - loss: 1.1588 - precision: 0.6108 - recall: 0.5708 - val\_accuracy: 0.8195 - val\_loss: 0.5499 - val\_precision: 0.8633 - val\_recall: 0.7639 - learning\_rate: 1.0000e-04  
Epoch 2/25

**274/274** 6s 24ms/step - accuracy: 0.7262 - loss: 0.7887 - precision: 0.7460 - recall: 0.7060 - val\_accuracy: 0.8745 - val\_loss: 0.3641 - val\_precision: 0.8825 - val\_recall: 0.8596 - learning\_rate: 1.0000e-04  
Epoch 3/25

**274/274** 7s 25ms/step - accuracy: 0.7789 - loss: 0.6360 - precision: 0.7951 - recall: 0.7609 - val\_accuracy: 0.9004 - val\_loss: 0.2842 - val\_precision: 0.9062 - val\_recall: 0.8875 - learning\_rate: 1.0000e-04  
Epoch 4/25

**274/274** 11s 27ms/step - accuracy: 0.8115 - loss: 0.5508 - precision: 0.8247 - recall: 0.7981 - val\_accuracy: 0.9004 - val\_loss: 0.3271 - val\_precision: 0.9041 - val\_recall: 0.8900 - learning\_rate: 1.0000e-04  
Epoch 5/25

**274/274** 7s 24ms/step - accuracy: 0.8357 - loss: 0.4786 - precision: 0.8450 - recall: 0.8235 - val\_accuracy: 0.9159 - val\_loss: 0.2373 - val\_precision: 0.9221 - val\_recall: 0.9114 - learning\_rate: 1.0000e-04  
Epoch 6/25

**274/274** 7s 25ms/step - accuracy: 0.8583 - loss: 0.4247 - precision: 0.8652 - recall: 0.8480 - val\_accuracy: 0.9172 - val\_loss: 0.2335 - val\_precision: 0.9216 - val\_recall: 0.9127 - learning\_rate: 1.0000e-04  
Epoch 7/25

**274/274** 7s 25ms/step - accuracy: 0.8679 - loss: 0.3899 - precision: 0.8761 - recall: 0.8577 - val\_accuracy: 0.9230 - val\_loss: 0.2233 - val\_precision: 0.9258 - val\_recall: 0.9204 - learning\_rate: 1.0000e-04  
Epoch 8/25

**274/274** 8s 28ms/step - accuracy: 0.8734 - loss: 0.3686 - precision: 0.8797 - recall: 0.8660 - val\_accuracy: 0.9295 - val\_loss: 0.2080 - val\_precision: 0.9316 - val\_recall: 0.9256 - learning\_rate: 1.0000e-04  
Epoch 9/25

**274/274** 7s 25ms/step - accuracy: 0.8795 - loss: 0.3461 - precision: 0.8851 - recall: 0.8733 - val\_accuracy: 0.9230 - val\_loss: 0.2014 - val\_precision: 0.9283 - val\_recall: 0.9211 - learning\_rate: 1.0000e-04  
Epoch 10/25

**274/274** 7s 25ms/step - accuracy: 0.8864 - loss: 0.3268 - precision: 0.8909 - recall: 0.8810 - val\_accuracy: 0.9386 - val\_loss: 0.1867 - val\_precision: 0.9401 - val\_recall: 0.9340 - learning\_rate: 1.0000e-04  
Epoch 11/25

**274/274** 7s 25ms/step - accuracy: 0.8864 - loss: 0.3118 - precision: 0.8920 - recall: 0.8809 - val\_accuracy: 0.9411 - val\_loss: 0.1850 - val\_precision: 0.9439 - val\_recall: 0.9366 - learning\_rate: 1.0000e-04  
Epoch 12/25

**274/274** 7s 24ms/step - accuracy: 0.8947 - loss: 0.2940 - precision: 0.8990 - recall: 0.8904 - val\_accuracy: 0.9411 - val\_loss: 0.1805 - val\_precision: 0.9421 - val\_recall: 0.9373 - learning\_rate: 1.0000e-04  
Epoch 13/25

**274/274** 7s 27ms/step - accuracy: 0.8972 - loss: 0.2907 - precision: 0.9007 - recall: 0.8919 - val\_accuracy: 0.9340 - val\_loss: 0.1800 - val\_precision: 0.9362 - val\_recall: 0.9308 - learning\_rate: 1.0000e-04  
Epoch 14/25

**274/274** 7s 25ms/step - accuracy: 0.9021 - loss: 0.2782 - precision: 0.9056 - recall: 0.8977 - val\_accuracy: 0.9360 - val\_loss: 0.1778 - val\_precision: 0.9389 - val\_recall: 0.9317 - learning\_rate: 1.0000e-04

```

al_precision: 0.9389 - val_recall: 0.9347 - learning_rate: 1.0000e-04
Epoch 15/25
274/274 7s 24ms/step - accuracy: 0.9048 - loss: 0.2607 - p
recision: 0.9075 - recall: 0.9005 - val_accuracy: 0.9437 - val_loss: 0.1661 - v
al_precision: 0.9460 - val_recall: 0.9398 - learning_rate: 1.0000e-04
Epoch 16/25
274/274 7s 26ms/step - accuracy: 0.9034 - loss: 0.2624 - p
recision: 0.9065 - recall: 0.9003 - val_accuracy: 0.9495 - val_loss: 0.1645 - v
al_precision: 0.9493 - val_recall: 0.9444 - learning_rate: 1.0000e-04
Epoch 17/25
274/274 7s 24ms/step - accuracy: 0.9039 - loss: 0.2611 - p
recision: 0.9076 - recall: 0.9010 - val_accuracy: 0.9198 - val_loss: 0.2099 - v
al_precision: 0.9226 - val_recall: 0.9172 - learning_rate: 1.0000e-04
Epoch 18/25
274/274 7s 26ms/step - accuracy: 0.9127 - loss: 0.2393 - p
recision: 0.9154 - recall: 0.9098 - val_accuracy: 0.9534 - val_loss: 0.1619 - v
al_precision: 0.9546 - val_recall: 0.9528 - learning_rate: 1.0000e-04
Epoch 19/25
274/274 7s 24ms/step - accuracy: 0.9145 - loss: 0.2265 - p
recision: 0.9177 - recall: 0.9112 - val_accuracy: 0.9463 - val_loss: 0.1546 - v
al_precision: 0.9481 - val_recall: 0.9444 - learning_rate: 1.0000e-04
Epoch 20/25
274/274 6s 24ms/step - accuracy: 0.9153 - loss: 0.2299 - p
recision: 0.9182 - recall: 0.9127 - val_accuracy: 0.9398 - val_loss: 0.1645 - v
al_precision: 0.9416 - val_recall: 0.9386 - learning_rate: 1.0000e-04
Epoch 21/25
274/274 6s 24ms/step - accuracy: 0.9176 - loss: 0.2214 - p
recision: 0.9206 - recall: 0.9148 - val_accuracy: 0.9314 - val_loss: 0.1770 - v
al_precision: 0.9320 - val_recall: 0.9301 - learning_rate: 1.0000e-04
Epoch 22/25
274/274 7s 24ms/step - accuracy: 0.9216 - loss: 0.2154 - p
recision: 0.9249 - recall: 0.9194 - val_accuracy: 0.9159 - val_loss: 0.2180 - v
al_precision: 0.9176 - val_recall: 0.9153 - learning_rate: 1.0000e-04
Epoch 23/25
273/274 0s 23ms/step - accuracy: 0.9247 - loss: 0.2037 - p
recision: 0.9270 - recall: 0.9227
Epoch 23: ReduceLROnPlateau reducing learning rate to 4.999999873689376e-05.
274/274 8s 28ms/step - accuracy: 0.9205 - loss: 0.2142 - p
recision: 0.9223 - recall: 0.9181 - val_accuracy: 0.9476 - val_loss: 0.1553 - v
al_precision: 0.9494 - val_recall: 0.9457 - learning_rate: 1.0000e-04
Epoch 24/25
274/274 7s 25ms/step - accuracy: 0.9302 - loss: 0.2016 - p
recision: 0.9321 - recall: 0.9282 - val_accuracy: 0.9534 - val_loss: 0.1549 - v
al_precision: 0.9553 - val_recall: 0.9534 - learning_rate: 5.0000e-05
Epoch 25/25
274/274 7s 24ms/step - accuracy: 0.9251 - loss: 0.2015 - p
recision: 0.9264 - recall: 0.9222 - val_accuracy: 0.9567 - val_loss: 0.1464 - v
al_precision: 0.9585 - val_recall: 0.9567 - learning_rate: 5.0000e-05
Restoring model weights from the end of the best epoch: 25.

```

```

In [ ]: # Evaluate VGG19
print("\nVGG19 - Test Set Evaluation:")
vgg19_test_loss, vgg19_test_acc, vgg19_test_prec, vgg19_test_recall = vgg19_mc
print(f"Test Accuracy: {vgg19_test_acc:.4f}")

```

```

print(f"Test Precision: {vgg19_test_prec:.4f}")
print(f"Test Recall: {vgg19_test_recall:.4f}")
print(f"Test F1-Score: {2 * (vgg19_test_prec * vgg19_test_recall) / (vgg19_te
test_results = {
    "model_name": "VGG19",
    "test_loss": float(vgg19_test_loss),
    "test_accuracy": float(vgg19_test_acc),
    "test_precision": float(vgg19_test_prec),
    "test_recall": float(vgg19_test_recall),
    "test_f1": float(2 * (vgg19_test_prec * vgg19_test_recall) / (vgg19_test_p
}

```

VGG19 - Test Set Evaluation:

**152/152** ————— **4s** 29ms/step - accuracy: 0.9464 - loss: 0.1515 - p  
recision: 0.9473 - recall: 0.9460  
Test Accuracy: 0.9464  
Test Precision: 0.9473  
Test Recall: 0.9460  
Test F1-Score: 0.9467

In [ ]: `save_model(vgg19_model, "VGG19", vgg19_history, test_results)`

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save\_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my\_model.keras')` or `keras.saving.save\_model(model, 'my\_model.keras')`.

## Model 2: InceptionV3

In [ ]: `def create_inception_model(input_shape=(192, 192, 3), num_classes=3):
 """
 Create InceptionV3 transfer learning model.
 """
 base_model = InceptionV3(
 weights='imagenet',
 include_top=False,
 input_shape=input_shape
 )

 base_model.trainable = False

 model = Sequential([
 base_model,
 GlobalAveragePooling2D(),
 Dense(512, activation='relu'),
 BatchNormalization(),
 Dropout(0.5),
 Dense(256, activation='relu'),
 BatchNormalization(),
 Dropout(0.5),
 Dense(num_classes, activation='softmax')
 ], name='InceptionV3_Transfer')`

```
    return model

# Create and compile InceptionV3
inception_model = create_inception_model()
inception_model.compile(
    optimizer=Adam(learning_rate=0.0001),
    loss='categorical_crossentropy',
    metrics=['accuracy',
              tf.keras.metrics.Precision(name='precision'),
              tf.keras.metrics.Recall(name='recall')])
)

print("InceptionV3 Model Created")
```

InceptionV3 Model Created

```
In [ ]: # Train InceptionV3
print("Training InceptionV3...")
inception_history = inception_model.fit(
    X_train_split, y_train_split,
    validation_data=(X_val, y_val),
    epochs=25,
    batch_size=32,
    callbacks=setup_callbacks(),
    verbose=1
)
```

Training InceptionV3...

Epoch 1/25

**274/274** 22s 50ms/step - accuracy: 0.5943 - loss: 1.1937 - precision: 0.6121 - recall: 0.5691 - val\_accuracy: 0.8292 - val\_loss: 0.5139 - val\_precision: 0.8507 - val\_recall: 0.7995 - learning\_rate: 1.0000e-04

Epoch 2/25

**274/274** 3s 12ms/step - accuracy: 0.7221 - loss: 0.8014 - precision: 0.7369 - recall: 0.6990 - val\_accuracy: 0.8396 - val\_loss: 0.4887 - val\_precision: 0.8542 - val\_recall: 0.8150 - learning\_rate: 1.0000e-04

Epoch 3/25

**274/274** 4s 14ms/step - accuracy: 0.7810 - loss: 0.6386 - precision: 0.7981 - recall: 0.7625 - val\_accuracy: 0.8829 - val\_loss: 0.3445 - val\_precision: 0.8866 - val\_recall: 0.8752 - learning\_rate: 1.0000e-04

Epoch 4/25

**274/274** 3s 11ms/step - accuracy: 0.8103 - loss: 0.5511 - precision: 0.8244 - recall: 0.7938 - val\_accuracy: 0.8829 - val\_loss: 0.3380 - val\_precision: 0.8890 - val\_recall: 0.8758 - learning\_rate: 1.0000e-04

Epoch 5/25

**274/274** 3s 11ms/step - accuracy: 0.8292 - loss: 0.5128 - precision: 0.8388 - recall: 0.8171 - val\_accuracy: 0.8849 - val\_loss: 0.3452 - val\_precision: 0.8877 - val\_recall: 0.8745 - learning\_rate: 1.0000e-04

Epoch 6/25

**274/274** 5s 12ms/step - accuracy: 0.8401 - loss: 0.4618 - precision: 0.8485 - recall: 0.8307 - val\_accuracy: 0.8972 - val\_loss: 0.3010 - val\_precision: 0.9020 - val\_recall: 0.8933 - learning\_rate: 1.0000e-04

Epoch 7/25

**274/274** 3s 11ms/step - accuracy: 0.8537 - loss: 0.4279 - precision: 0.8610 - recall: 0.8427 - val\_accuracy: 0.8972 - val\_loss: 0.2846 - val\_precision: 0.8992 - val\_recall: 0.8939 - learning\_rate: 1.0000e-04

Epoch 8/25

**274/274** 3s 11ms/step - accuracy: 0.8638 - loss: 0.3952 - precision: 0.8697 - recall: 0.8548 - val\_accuracy: 0.8959 - val\_loss: 0.2917 - val\_precision: 0.8981 - val\_recall: 0.8946 - learning\_rate: 1.0000e-04

Epoch 9/25

**274/274** 3s 12ms/step - accuracy: 0.8629 - loss: 0.3886 - precision: 0.8681 - recall: 0.8556 - val\_accuracy: 0.9062 - val\_loss: 0.2640 - val\_precision: 0.9073 - val\_recall: 0.9049 - learning\_rate: 1.0000e-04

Epoch 10/25

**274/274** 3s 11ms/step - accuracy: 0.8722 - loss: 0.3680 - precision: 0.8766 - recall: 0.8645 - val\_accuracy: 0.9056 - val\_loss: 0.2621 - val\_precision: 0.9082 - val\_recall: 0.9023 - learning\_rate: 1.0000e-04

Epoch 11/25

**274/274** 3s 11ms/step - accuracy: 0.8743 - loss: 0.3493 - precision: 0.8799 - recall: 0.8674 - val\_accuracy: 0.9056 - val\_loss: 0.2636 - val\_precision: 0.9077 - val\_recall: 0.9036 - learning\_rate: 1.0000e-04

Epoch 12/25

**274/274** 3s 11ms/step - accuracy: 0.8800 - loss: 0.3225 - precision: 0.8861 - recall: 0.8750 - val\_accuracy: 0.8991 - val\_loss: 0.2662 - val\_precision: 0.9008 - val\_recall: 0.8984 - learning\_rate: 1.0000e-04

Epoch 13/25

**274/274** 4s 13ms/step - accuracy: 0.8825 - loss: 0.3202 - precision: 0.8870 - recall: 0.8769 - val\_accuracy: 0.9030 - val\_loss: 0.2741 - val\_precision: 0.9074 - val\_recall: 0.9004 - learning\_rate: 1.0000e-04

Epoch 14/25

```
274/274 ----- 3s 11ms/step - accuracy: 0.8840 - loss: 0.3112 - precision: 0.8880 - recall: 0.8797 - val_accuracy: 0.9120 - val_loss: 0.2413 - val_precision: 0.9138 - val_recall: 0.9120 - learning_rate: 1.0000e-04
Epoch 15/25
274/274 ----- 3s 11ms/step - accuracy: 0.8856 - loss: 0.3123 - precision: 0.8892 - recall: 0.8822 - val_accuracy: 0.8946 - val_loss: 0.2841 - val_precision: 0.8963 - val_recall: 0.8946 - learning_rate: 1.0000e-04
Epoch 16/25
274/274 ----- 3s 11ms/step - accuracy: 0.8904 - loss: 0.2900 - precision: 0.8937 - recall: 0.8856 - val_accuracy: 0.9075 - val_loss: 0.2400 - val_precision: 0.9087 - val_recall: 0.9075 - learning_rate: 1.0000e-04
Epoch 17/25
274/274 ----- 3s 11ms/step - accuracy: 0.9005 - loss: 0.2706 - precision: 0.9029 - recall: 0.8968 - val_accuracy: 0.9107 - val_loss: 0.2317 - val_precision: 0.9113 - val_recall: 0.9101 - learning_rate: 1.0000e-04
Epoch 18/25
274/274 ----- 3s 11ms/step - accuracy: 0.8946 - loss: 0.2777 - precision: 0.8985 - recall: 0.8921 - val_accuracy: 0.9153 - val_loss: 0.2266 - val_precision: 0.9169 - val_recall: 0.9133 - learning_rate: 1.0000e-04
Epoch 19/25
274/274 ----- 3s 10ms/step - accuracy: 0.9013 - loss: 0.2600 - precision: 0.9041 - recall: 0.8983 - val_accuracy: 0.9017 - val_loss: 0.2520 - val_precision: 0.9028 - val_recall: 0.9010 - learning_rate: 1.0000e-04
Epoch 20/25
274/274 ----- 3s 10ms/step - accuracy: 0.9085 - loss: 0.2509 - precision: 0.9107 - recall: 0.9060 - val_accuracy: 0.9088 - val_loss: 0.2374 - val_precision: 0.9094 - val_recall: 0.9088 - learning_rate: 1.0000e-04
Epoch 21/25
274/274 ----- 3s 11ms/step - accuracy: 0.8989 - loss: 0.2580 - precision: 0.9020 - recall: 0.8963 - val_accuracy: 0.9094 - val_loss: 0.2485 - val_precision: 0.9100 - val_recall: 0.9088 - learning_rate: 1.0000e-04
Epoch 22/25
274/274 ----- 3s 10ms/step - accuracy: 0.9045 - loss: 0.2520 - precision: 0.9071 - recall: 0.9013 - val_accuracy: 0.9082 - val_loss: 0.2506 - val_precision: 0.9091 - val_recall: 0.9062 - learning_rate: 1.0000e-04
Epoch 23/25
272/274 ----- 0s 9ms/step - accuracy: 0.9099 - loss: 0.2476 - precision: 0.9121 - recall: 0.9077
Epoch 23: ReduceLROnPlateau reducing learning rate to 4.999999873689376e-05.
274/274 ----- 3s 11ms/step - accuracy: 0.9125 - loss: 0.2395 - precision: 0.9149 - recall: 0.9106 - val_accuracy: 0.9146 - val_loss: 0.2301 - val_precision: 0.9157 - val_recall: 0.9140 - learning_rate: 1.0000e-04
Epoch 24/25
274/274 ----- 4s 15ms/step - accuracy: 0.9133 - loss: 0.2308 - precision: 0.9150 - recall: 0.9112 - val_accuracy: 0.9133 - val_loss: 0.2245 - val_precision: 0.9139 - val_recall: 0.9133 - learning_rate: 5.0000e-05
Epoch 25/25
274/274 ----- 3s 11ms/step - accuracy: 0.9127 - loss: 0.2301 - precision: 0.9148 - recall: 0.9112 - val_accuracy: 0.9133 - val_loss: 0.2251 - val_precision: 0.9133 - val_recall: 0.9133 - learning_rate: 5.0000e-05
Restoring model weights from the end of the best epoch: 18.
```

```
In [ ]: # Evaluate InceptionV3
print("\nInceptionV3 - Test Set Evaluation:")
```

```

inception_test_loss, inception_test_acc, inception_test_prec, inception_test_r
print(f"Test Accuracy: {inception_test_acc:.4f}")
print(f"Test Precision: {inception_test_prec:.4f}")
print(f"Test Recall: {inception_test_recall:.4f}")
print(f"Test F1-Score: {2 * (inception_test_prec * inception_test_recall) / (inception_test_precision * inception_test_recall)}")

test_results = {
    "model_name": "InceptionV3",
    "test_loss": float(inception_test_loss),
    "test_accuracy": float(inception_test_acc),
    "test_precision": float(inception_test_prec),
    "test_recall": float(inception_test_recall),
    "test_f1": float(2 * (inception_test_prec * inception_test_recall) / (inception_test_precision * inception_test_recall))
}

```

InceptionV3 - Test Set Evaluation:  
**152/152** ————— 4s 28ms/step - accuracy: 0.9060 - loss: 0.2511 - precision: 0.9074 - recall: 0.9035  
Test Accuracy: 0.9060  
Test Precision: 0.9074  
Test Recall: 0.9035  
Test F1-Score: 0.9054

In [ ]: `save_model(inception_model, "InceptionV3", inception_history, test_results)`

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save\_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my\_model.keras')` or `keras.saving.save\_model(model, 'my\_model.keras')`.

### Model 3: mobilenetv2

In [ ]: `from tensorflow.keras.layers import Input, Conv2D, Rescaling, GlobalAveragePooling2D
from tensorflow.keras.applications import MobileNetV2

def create_mobilenetv2_model(input_shape=(192,192,1), num_classes=3):
 inputs = Input(shape=input_shape)
 x = inputs

 x = Rescaling(1./127.5, offset=-1)(x)

 base = MobileNetV2(weights='imagenet', include_top=False, input_shape=(192,192,3))
 base.trainable = False

 x = base(x)

 x = GlobalAveragePooling2D()(x)
 x = Dense(256, activation='relu')(x)
 x = BatchNormalization()(x)
 x = Dropout(0.5)(x)
 outputs = Dense(num_classes, activation='softmax')(x)

 return Model(inputs, outputs, name='MobileNetV2_Transfer')`

```

mobilenet_model = create_mobilenetv2_model()
mobilenet_model.compile(
    optimizer=Adam(1e-4),
    loss='categorical_crossentropy',
    metrics=['accuracy',
              tf.keras.metrics.Precision(name='precision'),
              tf.keras.metrics.Recall(name='recall')]
)
mobilenet_model.summary()

```

**Model: "MobileNetV2\_Transfer"**

Layer (type)	Output Shape	Param #
input_layer_7 (InputLayer)	(None, 192, 192, 1)	0
rescaling (Rescaling)	(None, 192, 192, 1)	0
mobilenetv2_1.00_192 (Functional)	(None, 6, 6, 1280)	2,257,984
global_average_pooling2d_3 (GlobalAveragePooling2D)	(None, 1280)	0
dense_12 (Dense)	(None, 256)	327,936
batch_normalization_108 (BatchNormalization)	(None, 256)	1,024
dropout_11 (Dropout)	(None, 256)	0
dense_13 (Dense)	(None, 3)	771

**Total params:** 2,587,715 (9.87 MB)

**Trainable params:** 329,219 (1.26 MB)

**Non-trainable params:** 2,258,496 (8.62 MB)

```

In [ ]: # Train EfficientNet
print("Training mobilenet..")
mobilenet_history = mobilenet_model.fit(
    X_train_split, y_train_split,
    validation_data=(X_val, y_val),
    epochs=25,
    batch_size=32,
    callbacks=setup_callbacks(),
    verbose=1
)

```

Training mobilenet..

Epoch 1/25

```
2025-10-30 17:57:22.570788: I external/local_xla/xla/service/gpu/autotuning/do
t_search_space.cc:208] All configs were filtered out because none of them suffi
ciently match the hints. Maybe the hints set does not contain a good representa
tive set of valid configs? Working around this by using the full hints set inst
ead.
2025-10-30 17:57:22.820174: I external/local_xla/xla/stream_executor/cuda/subpr
ocess_compilation.cc:346] ptxas warning : Registers are spilled to local memory
in function 'gemm_fusion_dot_2320', 4 bytes spill stores, 4 bytes spill loads

2025-10-30 17:57:22.913290: I external/local_xla/xla/stream_executor/cuda/subpr
ocess_compilation.cc:346] ptxas warning : Registers are spilled to local memory
in function 'gemm_fusion_dot_2320', 520 bytes spill stores, 520 bytes spill loa
ds

2025-10-30 17:57:26.479377: E external/local_xla/xla/stream_executor/cuda/cud
a_timer.cc:86] Delay kernel timed out: measured time has sub-optimal accuracy.
There may be a missing warmup execution, please investigate in Nsight Systems.
2025-10-30 17:57:26.629087: E external/local_xla/xla/stream_executor/cuda/cud
a_timer.cc:86] Delay kernel timed out: measured time has sub-optimal accuracy.
There may be a missing warmup execution, please investigate in Nsight Systems.
273/274 0s 6ms/step - accuracy: 0.5262 - loss: 1.0268 - pr
ecision: 0.5787 - recall: 0.3585
```

```
2025-10-30 17:57:33.285511: E external/local_xla/xla/stream_executor/cuda/cud
a_timer.cc:86] Delay kernel timed out: measured time has sub-optimal accuracy.
There may be a missing warmup execution, please investigate in Nsight Systems.
2025-10-30 17:57:33.450630: E external/local_xla/xla/stream_executor/cuda/cud
a_timer.cc:86] Delay kernel timed out: measured time has sub-optimal accuracy.
There may be a missing warmup execution, please investigate in Nsight Systems.
2025-10-30 17:57:33.602377: E external/local_xla/xla/stream_executor/cuda/cud
a_timer.cc:86] Delay kernel timed out: measured time has sub-optimal accuracy.
There may be a missing warmup execution, please investigate in Nsight Systems.
2025-10-30 17:57:33.777732: E external/local_xla/xla/stream_executor/cuda/cud
a_timer.cc:86] Delay kernel timed out: measured time has sub-optimal accuracy.
There may be a missing warmup execution, please investigate in Nsight Systems.
2025-10-30 17:57:33.925207: E external/local_xla/xla/stream_executor/cuda/cud
a_timer.cc:86] Delay kernel timed out: measured time has sub-optimal accuracy.
There may be a missing warmup execution, please investigate in Nsight Systems.
2025-10-30 17:57:34.068251: E external/local_xla/xla/service/slow_operation_ala
rm.cc:73] Trying algorithm eng3{k11=0} for conv (f32[22,960,6,6]{3,2,1,0}, u
8[0]{0}) custom-call(f32[22,960,6,6]{3,2,1,0}, f32[960,1,3,3]{3,2,1,0}), windo
w={size=3x3 pad=1_1x1_1}, dim_labels=bf01_oi01->bf01, feature_group_count=960,
custom_call_target="__cudnn$convForward", backend_config={"operation_queue_i
d":"0","wait_on_operation_queues":[],"cudnn_conv_backend_config":{"activation_m
ode":"kNone","conv_result_scale":1,"side_input_scale":0,"leakyrelu_alpha":0},"f
orce_earliest_schedule":false,"reification_cost":[]}) is taking a while...
2025-10-30 17:57:34.087643: E external/local_xla/xla/stream_executor/cuda/cud
a_timer.cc:86] Delay kernel timed out: measured time has sub-optimal accuracy.
There may be a missing warmup execution, please investigate in Nsight Systems.
2025-10-30 17:57:34.252589: E external/local_xla/xla/stream_executor/cuda/cud
a_timer.cc:86] Delay kernel timed out: measured time has sub-optimal accuracy.
There may be a missing warmup execution, please investigate in Nsight Systems.
2025-10-30 17:57:34.400752: E external/local_xla/xla/stream_executor/cuda/cud
a_timer.cc:86] Delay kernel timed out: measured time has sub-optimal accuracy.
There may be a missing warmup execution, please investigate in Nsight Systems.
2025-10-30 17:57:34.550431: E external/local_xla/xla/stream_executor/cuda/cud
a_timer.cc:86] Delay kernel timed out: measured time has sub-optimal accuracy.
There may be a missing warmup execution, please investigate in Nsight Systems.
2025-10-30 17:57:34.692959: E external/local_xla/xla/stream_executor/cuda/cud
a_timer.cc:86] Delay kernel timed out: measured time has sub-optimal accuracy.
There may be a missing warmup execution, please investigate in Nsight Systems.
2025-10-30 17:57:34.694093: E external/local_xla/xla/service/slow_operation_ala
rm.cc:140] The operation took 1.636592636s
Trying algorithm eng3{k11=0} for conv (f32[22,960,6,6]{3,2,1,0}, u8[0]{0}) cust
om-call(f32[22,960,6,6]{3,2,1,0}, f32[960,1,3,3]{3,2,1,0}), window={size=3x3 pa
d=1_1x1_1}, dim_labels=bf01_oi01->bf01, feature_group_count=960, custom_call_ta
rget="__cudnn$convForward", backend_config={"operation_queue_id":"0","wait_on_o
peration_queues":[],"cudnn_conv_backend_config":{"activation_mode":"kNone","con
v_result_scale":1,"side_input_scale":0,"leakyrelu_alpha":0},"force_earliest_sch
edule":false,"reification_cost":[]}) is taking a while...
```

**274/274** **0s** 29ms/step - accuracy: 0.5265 - loss: 1.0264 - p
recision: 0.5790 - recall: 0.3589

2025-10-30 17:57:41.702925: E external/local\_xla/xla/stream\_executor/cuda/cuda\_timer.cc:86] Delay kernel timed out: measured time has sub-optimal accuracy. There may be a missing warmup execution, please investigate in Nsight Systems.

2025-10-30 17:57:41.853981: E external/local\_xla/xla/stream\_executor/cuda/cuda\_timer.cc:86] Delay kernel timed out: measured time has sub-optimal accuracy. There may be a missing warmup execution, please investigate in Nsight Systems.

2025-10-30 17:57:42.000153: E external/local\_xla/xla/stream\_executor/cuda/cuda\_timer.cc:86] Delay kernel timed out: measured time has sub-optimal accuracy. There may be a missing warmup execution, please investigate in Nsight Systems.

**274/274** 24s 56ms/step - accuracy: 0.6037 - loss: 0.9017 - precision: 0.6730 - recall: 0.4551 - val\_accuracy: 0.2387 - val\_loss: 1.4298 - val\_precision: 0.2390 - val\_recall: 0.2387 - learning\_rate: 1.0000e-04  
Epoch 2/25

**274/274** 2s 7ms/step - accuracy: 0.7555 - loss: 0.6292 - precision: 0.8031 - recall: 0.6577 - val\_accuracy: 0.0977 - val\_loss: 1.6812 - val\_precision: 0.0964 - val\_recall: 0.0944 - learning\_rate: 1.0000e-04  
Epoch 3/25

**274/274** 2s 7ms/step - accuracy: 0.7921 - loss: 0.5208 - precision: 0.8199 - recall: 0.7383 - val\_accuracy: 0.7555 - val\_loss: 0.5127 - val\_precision: 0.7818 - val\_recall: 0.7322 - learning\_rate: 1.0000e-04  
Epoch 4/25

**274/274** 2s 7ms/step - accuracy: 0.8051 - loss: 0.4748 - precision: 0.8264 - recall: 0.7711 - val\_accuracy: 0.7380 - val\_loss: 0.6441 - val\_precision: 0.7731 - val\_recall: 0.7251 - learning\_rate: 1.0000e-04  
Epoch 5/25

**274/274** 3s 11ms/step - accuracy: 0.8114 - loss: 0.4493 - precision: 0.8293 - recall: 0.7848 - val\_accuracy: 0.2451 - val\_loss: 4.0727 - val\_precision: 0.2451 - val\_recall: 0.2451 - learning\_rate: 1.0000e-04  
Epoch 6/25

**274/274** 2s 7ms/step - accuracy: 0.8166 - loss: 0.4349 - precision: 0.8313 - recall: 0.7948 - val\_accuracy: 0.7426 - val\_loss: 0.6084 - val\_precision: 0.7463 - val\_recall: 0.7419 - learning\_rate: 1.0000e-04  
Epoch 7/25

**274/274** 2s 7ms/step - accuracy: 0.8224 - loss: 0.4202 - precision: 0.8360 - recall: 0.8005 - val\_accuracy: 0.6934 - val\_loss: 2.2518 - val\_precision: 0.6939 - val\_recall: 0.6934 - learning\_rate: 1.0000e-04  
Epoch 8/25

**274/274** 2s 7ms/step - accuracy: 0.8285 - loss: 0.4082 - precision: 0.8406 - recall: 0.8091 - val\_accuracy: 0.8512 - val\_loss: 0.3745 - val\_precision: 0.8687 - val\_recall: 0.8305 - learning\_rate: 1.0000e-04  
Epoch 9/25

**274/274** 2s 7ms/step - accuracy: 0.8340 - loss: 0.3998 - precision: 0.8475 - recall: 0.8170 - val\_accuracy: 0.2536 - val\_loss: 3.6650 - val\_precision: 0.2536 - val\_recall: 0.2536 - learning\_rate: 1.0000e-04  
Epoch 10/25

**274/274** 2s 7ms/step - accuracy: 0.8381 - loss: 0.3905 - precision: 0.8479 - recall: 0.8202 - val\_accuracy: 0.6630 - val\_loss: 0.7953 - val\_precision: 0.6883 - val\_recall: 0.6255 - learning\_rate: 1.0000e-04  
Epoch 11/25

**274/274** 2s 7ms/step - accuracy: 0.8385 - loss: 0.3839 - precision: 0.8483 - recall: 0.8230 - val\_accuracy: 0.7594 - val\_loss: 0.5900 - val\_precision: 0.7656 - val\_recall: 0.7523 - learning\_rate: 1.0000e-04  
Epoch 12/25

**274/274** 2s 7ms/step - accuracy: 0.8443 - loss: 0.3754 - precision: 0.8529 - recall: 0.8309 - val\_accuracy: 0.6281 - val\_loss: 0.8250 - val\_precision: 0.6281 - val\_recall: 0.6216 - learning\_rate: 1.0000e-04  
Epoch 13/25

**270/274** 0s 6ms/step - accuracy: 0.8416 - loss: 0.3780 - precision: 0.8514 - recall: 0.8270  
Epoch 13: ReduceLROnPlateau reducing learning rate to 4.99999873689376e-05.  
**274/274** 2s 7ms/step - accuracy: 0.8415 - loss: 0.3730 - precision: 0.8508 - recall: 0.8293 - val\_accuracy: 0.8234 - val\_loss: 0.4062 - val\_precision: 0.8301 - val\_recall: 0.8124 - learning\_rate: 1.0000e-04

```
Epoch 14/25
274/274 2s 7ms/step - accuracy: 0.8470 - loss: 0.3635 - precision: 0.8562 - recall: 0.8371 - val_accuracy: 0.7167 - val_loss: 0.6962 - val_precision: 0.7197 - val_recall: 0.7141 - learning_rate: 5.0000e-05
Epoch 15/25
274/274 2s 7ms/step - accuracy: 0.8511 - loss: 0.3583 - precision: 0.8593 - recall: 0.8403 - val_accuracy: 0.7393 - val_loss: 0.9062 - val_precision: 0.7411 - val_recall: 0.7387 - learning_rate: 5.0000e-05
Epoch 16/25
274/274 2s 7ms/step - accuracy: 0.8519 - loss: 0.3494 - precision: 0.8608 - recall: 0.8405 - val_accuracy: 0.7477 - val_loss: 0.7956 - val_precision: 0.7503 - val_recall: 0.7464 - learning_rate: 5.0000e-05
Epoch 16: early stopping
Restoring model weights from the end of the best epoch: 8.
```

```
In [ ]: # Evaluate EfficientNet
print("\nmobilenet - Test Set Evaluation:")
mobilenet_test_loss, mobilenet_test_acc, mobilenet_test_prec, mobilenet_test_rec
print(f"Test Accuracy: {mobilenet_test_acc:.4f}")
print(f"Test Precision: {mobilenet_test_prec:.4f}")
print(f"Test Recall: {mobilenet_test_recall:.4f}")
print(f"Test F1-Score: {2 * (mobilenet_test_prec * mobilenet_test_recall) / (mobilenet_test_prec + mobilenet_test_recall)}")

test_results = {
    "model_name": "mobilenet",
    "test_loss": float(mobilenet_test_loss),
    "test_accuracy": float(mobilenet_test_acc),
    "test_precision": float(mobilenet_test_prec),
    "test_recall": float(mobilenet_test_recall),
    "test_f1": float(2 * (mobilenet_test_prec * mobilenet_test_recall) / (mobilenet_test_prec + mobilenet_test_recall))
}
```

```
mobilenet - Test Set Evaluation:
149/152 0s 6ms/step - accuracy: 0.8521 - loss: 0.3739 - precision: 0.8648 - recall: 0.8341
2025-10-30 17:58:43.395389: E external/local_xla/xla/stream_executor/cuda/cuda_timer.cc:86] Delay kernel timed out: measured time has sub-optimal accuracy. There may be a missing warmup execution, please investigate in Nsight Systems.
2025-10-30 17:58:43.536620: E external/local_xla/xla/stream_executor/cuda/cuda_timer.cc:86] Delay kernel timed out: measured time has sub-optimal accuracy. There may be a missing warmup execution, please investigate in Nsight Systems.
152/152 5s 35ms/step - accuracy: 0.8490 - loss: 0.3835 - precision: 0.8596 - recall: 0.8292
Test Accuracy: 0.8490
Test Precision: 0.8596
Test Recall: 0.8292
Test F1-Score: 0.8441
```

```
In [ ]: save_model(mobilenet_model, "mobilenet", mobilenet_history, test_results)
```

```
WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.
```

## 5. Performance Comparison

```
In [ ]: # Create comparison dataframe
models_summary = {
    'Model': ['Baseline CNN', 'ResNet50', 'VGG19', 'InceptionV3', 'mobilenet'],
    'Test Accuracy': [
        baseline_test_acc,
        resnet_test_acc,
        vgg19_test_acc,
        inception_test_acc,
        mobilenet_test_acc
    ],
    'Test Precision': [
        baseline_test_prec,
        resnet_test_prec,
        vgg19_test_prec,
        inception_test_prec,
        mobilenet_test_prec
    ],
    'Test Recall': [
        baseline_test_recall,
        resnet_test_recall,
        vgg19_test_recall,
        inception_test_recall,
        mobilenet_test_recall
    ],
    'Optimizer': ['Adam', 'Adam', 'Adam', 'Adam', 'Adam'],
    'Learning Rate': [0.001, 0.0001, 0.0001, 0.0001, 0.0001],
    'Batch Size': [32, 32, 32, 32, 32],
    'Max Epochs': [30, 35, 25, 25, 25]
}

comparison_df = pd.DataFrame(models_summary)
comparison_df['Test F1-Score'] = 2 * (comparison_df['Test Precision'] * comparison_df['Test Recall']) / (comparison_df['Test Precision'] + comparison_df['Test Recall'])

# Format percentages
for col in ['Test Accuracy', 'Test Precision', 'Test Recall', 'Test F1-Score']:
    comparison_df[col] = comparison_df[col].apply(lambda x: f"{x:.4f}")

print("\n" + "="*100)
print("MODEL PERFORMANCE COMPARISON")
print("="*100)
print(comparison_df.to_string(index=False))
print("="*100)

comparison_df.to_csv("out/comparison_df.csv")
```

Model	Batch Size	Test Max Epochs	Accuracy	Precision	Recall	Optimizer	Learning Rate
			Test F1-Score				
Baseline CNN	32	30	0.9810	0.9812	0.9806	Adam	0.00010
ResNet50	32	35	0.9809	0.9338	0.9357	Adam	0.0001
VGG19	32	25	0.9341	0.9464	0.9473	Adam	0.0001
InceptionV3	32	25	0.9467	0.9060	0.9074	Adam	0.0001
mobilenet	32	25	0.9054	0.8490	0.8596	Adam	0.0001
			0.8441		0.8292		

```
In [ ]: # Identify best model
comparison_df_numeric = pd.DataFrame(models_summary)
comparison_df_numeric['Test F1-Score'] = 2 * (comparison_df_numeric['Test Precision'] - comparison_df_numeric['Test Recall'])

best_model_idx = comparison_df_numeric['Test Accuracy'].idxmax()
best_model_name = comparison_df_numeric.loc[best_model_idx, 'Model']

print(f"\nBEST PERFORMING MODEL: {best_model_name}")
print(f"    Test Accuracy: {comparison_df_numeric.loc[best_model_idx, 'Test Accuracy']}")
print(f"    Test F1-Score: {comparison_df_numeric.loc[best_model_idx, 'Test F1-Score']}")
```

BEST PERFORMING MODEL: Baseline CNN  
 Test Accuracy: 0.9810  
 Test F1-Score: 0.9809

## Training Curves Comparison

```
In [ ]: # Plot training and validation curves for all models
fig, axes = plt.subplots(2, 3, figsize=(18, 10))
fig.suptitle('Training and Validation Performance Across Models', fontsize=16)

models_data = [
    ('Baseline CNN', baseline_history),
    ('ResNet50', resnet_history),
    ('VGG19', vgg19_history),
    ('InceptionV3', inception_history),
    ('mobilenet', mobilenet_history)
]

for idx, (name, history) in enumerate(models_data):
    row = idx // 3
    col = idx % 3
```

```

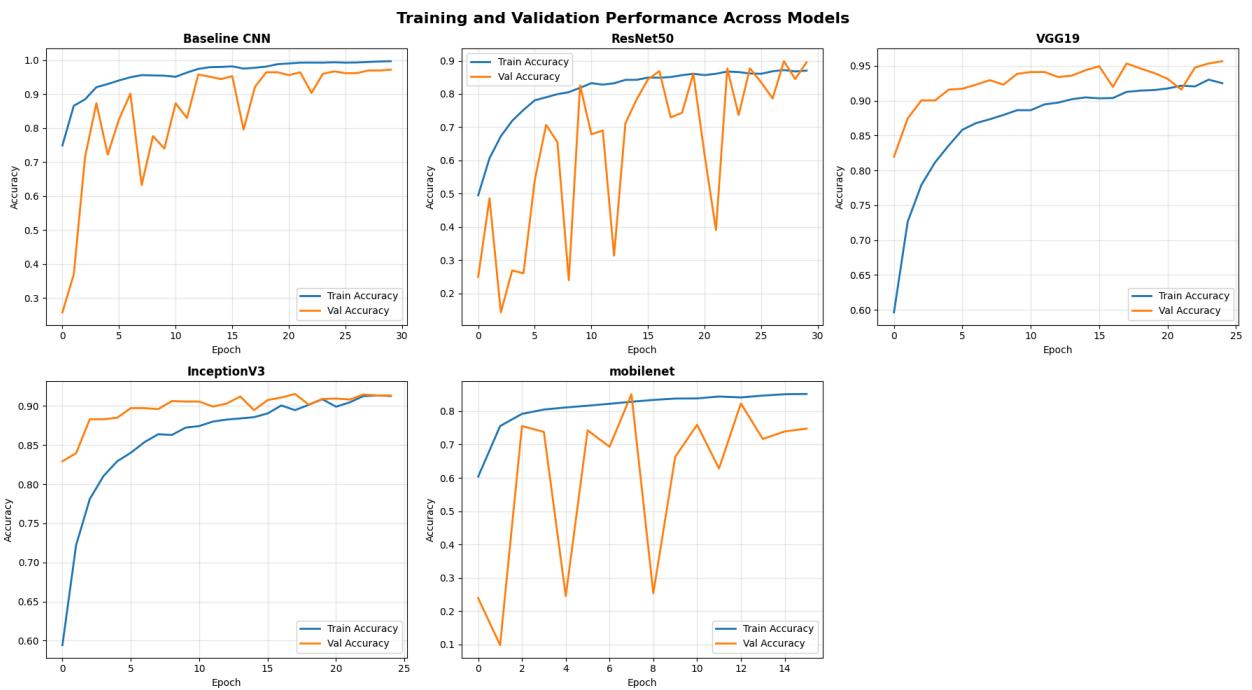
ax = axes[row, col]

# Plot accuracy
ax.plot(history.history['accuracy'], label='Train Accuracy', linewidth=2)
ax.plot(history.history['val_accuracy'], label='Val Accuracy', linewidth=2)
ax.set_title(name, fontweight='bold', fontsize=12)
ax.set_xlabel('Epoch')
ax.set_ylabel('Accuracy')
ax.legend()
ax.grid(True, alpha=0.3)

# Remove empty subplot
fig.delaxes(axes[1, 2])

plt.tight_layout()
plt.show()

```



```

In [ ]: # Plot loss curves
fig, axes = plt.subplots(2, 3, figsize=(18, 10))
fig.suptitle('Training and Validation Loss Across Models', fontsize=16, fontweight='bold')

for idx, (name, history) in enumerate(models_data):
    row = idx // 3
    col = idx % 3
    ax = axes[row, col]

    # Plot loss
    ax.plot(history.history['loss'], label='Train Loss', linewidth=2)
    ax.plot(history.history['val_loss'], label='Val Loss', linewidth=2)
    ax.set_title(name, fontweight='bold', fontsize=12)
    ax.set_xlabel('Epoch')
    ax.set_ylabel('Loss')

```

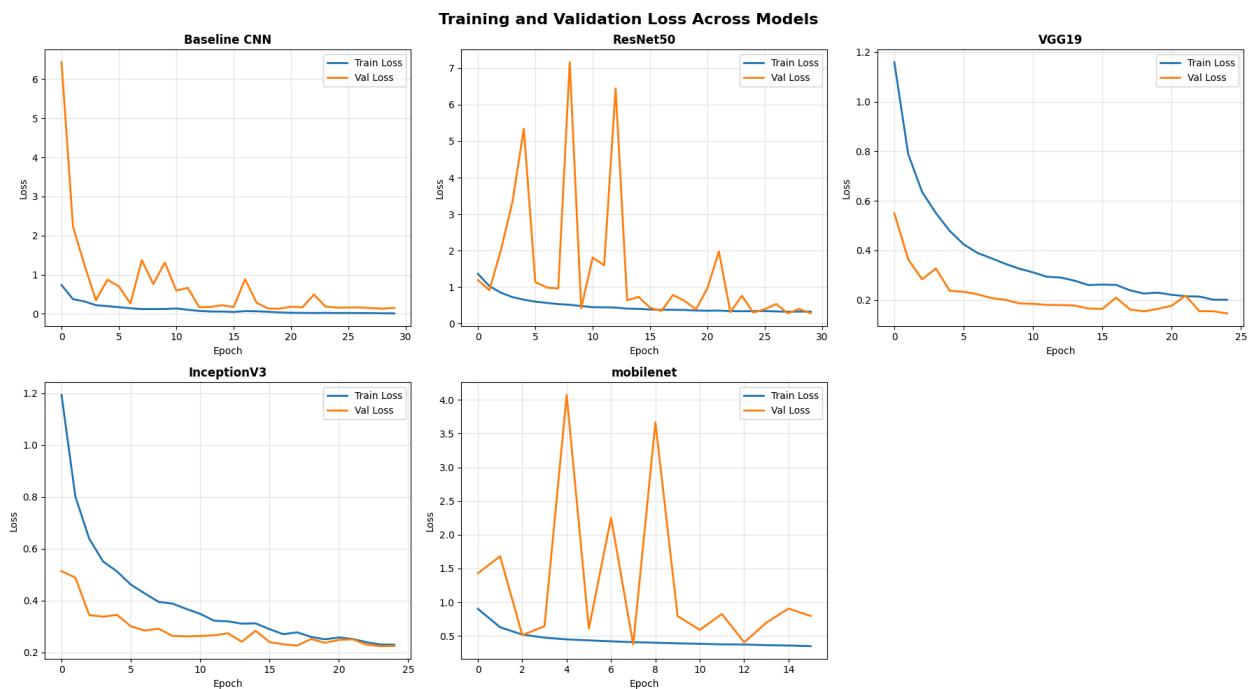
```

    ax.legend()
    ax.grid(True, alpha=0.3)

# Remove empty subplot
fig.delaxes(axes[1, 2])

plt.tight_layout()
plt.show()

```



## 6. Augmentation

### Augmentation Techniques Applied:

1. **Rotation:**  $\pm 15$  degrees
2. **Width/Height Shifts:**  $\pm 10\%$
3. **Horizontal Flip:** Random
4. **Zoom:**  $\pm 10\%$

**Note:** These augmentations were already applied during data balancing (Part 1). Here we analyze their impact by training the best model on both augmented and non-augmented data.

```

In [ ]: for name in dir():
    if not name.startswith('_'):
        del globals()[name]

import gc
gc.collect()

```

```
Out[ ]: 0
```

```
In [ ]: import numpy as np
import pandas as pd
from matplotlib import pyplot as plt
import tensorflow as tf
import pickle
from sklearn.model_selection import train_test_split

from tensorflow.python.keras.utils import np_utils
from tensorflow.keras.models import Sequential, Model

from sklearn.metrics import classification_report, confusion_matrix, accuracy_

from tensorflow.keras.layers import Dense, Dropout, Flatten, Activation, Batch
from tensorflow.keras.layers import Conv2D, MaxPooling2D

from tensorflow.keras.callbacks import EarlyStopping, ReduceLROnPlateau, Model
# from tensorflow.python.keras.layers.convolutional import Conv2D, MaxPooling2
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.optimizers import Adam, SGD, Adagrad, Adadelta, RMSprop
```

```
2025-10-30 18:11:20.553448: I tensorflow/core/util/port.cc:153] oneDNN custom operations are on. You may see slightly different numerical results due to floating-point round-off errors from different computation orders. To turn them off, set the environment variable `TF_ENABLE_ONEDNN_OPTS=0`.
2025-10-30 18:11:20.9555972: I tensorflow/core/platform/cpu_feature_guard.cc:210] This TensorFlow binary is optimized to use available CPU instructions in performance-critical operations.
To enable the following instructions: AVX2 AVX_VNNI FMA, in other operations, rebuild TensorFlow with the appropriate compiler flags.
2025-10-30 18:11:28.356135: I tensorflow/core/util/port.cc:153] oneDNN custom operations are on. You may see slightly different numerical results due to floating-point round-off errors from different computation orders. To turn them off, set the environment variable `TF_ENABLE_ONEDNN_OPTS=0`.
```

```
In [ ]: with open('data/X_test.pkl', 'rb') as file:
    X_test = pickle.load(file)

    with open('data/y_test.pkl', 'rb') as file:
        y_test = pickle.load(file)
```

```
In [ ]: # Load augmented training data for comparison
print("Loading augmented training data...")

with open('data/X_train_balanced.pkl', 'rb') as file:
    X_train_balanced = pickle.load(file)

with open('data/y_train_balanced.pkl', 'rb') as file:
    y_train_balanced = pickle.load(file)

print(f"Augmented training data shape: {X_train_balanced.shape}")
```

```
Loading augmented training data...
Augmented training data shape: (20713, 192, 192, 3)
```

```
In [ ]: # Create validation split from original data
X_train_balanced_split, X_val_balanced, y_train_balanced_split, y_val_balanced
    X_train_balanced, y_train_balanced,
    test_size=0.15,
    random_state=42,
    stratify=y_train_balanced
)

print(f"Balanced training split: {X_train_balanced_split.shape}")
print(f"Balanced validation split: {X_val_balanced.shape}")
```

```
Balanced training split: (17606, 192, 192, 3)
Balanced validation split: (3107, 192, 192, 3)
```

```
In [ ]: def create_baseline_cnn(input_shape=(192, 192, 3), num_classes=3):
    """
    Create a baseline CNN model.

    Architecture:
    - Conv Block 1: 32 filters, 3x3 kernel
    - Conv Block 2: 64 filters, 3x3 kernel
    - Conv Block 3: 128 filters, 3x3 kernel
    - Dense Layer: 256 units
    - Output Layer: 3 units (softmax)
    """

    model = Sequential([
        # First Conv Block
        Conv2D(32, (3, 3), activation='relu', padding='same', input_shape=input_shape),
        BatchNormalization(),
        Conv2D(32, (3, 3), activation='relu', padding='same'),
        BatchNormalization(),
        MaxPooling2D(pool_size=(2, 2)),
        Dropout(0.25),

        # Second Conv Block
        Conv2D(64, (3, 3), activation='relu', padding='same'),
        BatchNormalization(),
        Conv2D(64, (3, 3), activation='relu', padding='same'),
        BatchNormalization(),
        MaxPooling2D(pool_size=(2, 2)),
        Dropout(0.25),

        # Third Conv Block
        Conv2D(128, (3, 3), activation='relu', padding='same'),
        BatchNormalization(),
        Conv2D(128, (3, 3), activation='relu', padding='same'),
        BatchNormalization(),
        MaxPooling2D(pool_size=(2, 2)),
        Dropout(0.25),

        # Dense Layers
        Dense(256, activation='relu'),
        Dropout(0.5),
        Dense(num_classes, activation='softmax')
    ])

    return model
```

```

        Flatten(),
        Dense(256, activation='relu'),
        BatchNormalization(),
        Dropout(0.5),
        Dense(128, activation='relu'),
        BatchNormalization(),
        Dropout(0.5),

        # Output Layer
        Dense(num_classes, activation='softmax')
    ], name='Baseline_CNN')

    return model

```

```

In [ ]: # Setup callbacks
def setup_callbacks():
    early_stopping = EarlyStopping(
        monitor='val_accuracy',
        mode='max',
        min_delta=1e-3,
        patience=8,
        restore_best_weights=True,
        verbose=1
    )

    reduce_lr = ReduceLROnPlateau(
        monitor='val_accuracy',
        factor=0.5,
        patience=5,
        min_lr=1e-7,
        verbose=1
    )
    return [early_stopping, reduce_lr]

```

```

In [ ]: def save_model(model, model_name, model_history, test_results):
    import json
    import os
    os.makedirs("out", exist_ok=True)
    os.makedirs("out/model", exist_ok=True)
    os.makedirs("out/model_history", exist_ok=True)
    os.makedirs("out/test_result", exist_ok=True)

    model.save(f"out/model/{model_name}_model.h5")
    history_dict = getattr(model_history, "history", model_history)
    with open(f"out/model_history/{model_name}_history.json", "w") as f:
        json.dump(history_dict, f, indent=2)

    with open(f"out/test_result/{model_name}_test_result.json", "w") as f:
        json.dump(test_results, f, indent=2)

```

```

In [ ]: # Train best model on augmented data
print(f"\nTraining BASELINE CNN on BALANCED data (with augmentation)...")

```

```

# Recreate the best model based on comparison results

model_aug = create_baseline_cnn()

model_aug.compile(
    optimizer=Adam(learning_rate=0.001),
    loss='categorical_crossentropy',
    metrics=['accuracy',
              tf.keras.metrics.Precision(name='precision'),
              tf.keras.metrics.Recall(name='recall')])
)

history_aug = model_aug.fit(
    X_train_balanced_split, y_train_balanced_split,
    validation_data=(X_val_balanced, y_val_balanced),
    epochs=25,
    batch_size=32,
    callbacks=setup_callbacks(),
    verbose=1
)

```

Training BASELINE CNN on BALANCED data (with augmentation)...

```

/mnt/f/BaiduSyncdisk/1_Classes/3_25Fall/1_ADS/HW/HW_3A_Due_2025.10.12/.venv/lib/python3.12/site-packages/keras/src/layers/convolutional/base_conv.py:113: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.
    super().__init__(activity_regularizer=activity_regularizer, **kwargs)
WARNING: All log messages before absl::InitializeLog() is called are written to STDERR
I0000 00:00:1761862391.069649 54073 gpu_device.cc:2020] Created device /job:localhost/replica:0/task:0/device:GPU:0 with 21458 MB memory: -> device: 0, name: NVIDIA GeForce RTX 4090, pci bus id: 0000:01:00.0, compute capability: 8.9
Epoch 1/25

```

```
2025-10-30 18:13:23.601722: I external/local_xla/xla/service/service.cc:163] XL  
A service 0x78f744008e20 initialized for platform CUDA (this does not guarantee  
that XLA will be used). Devices:  
2025-10-30 18:13:23.601744: I external/local_xla/xla/service/service.cc:171]  
StreamExecutor device (0): NVIDIA GeForce RTX 4090, Compute Capability 8.9  
2025-10-30 18:13:23.760540: I tensorflow/compiler/mlir/tensorflow/utils/dump_ml  
ir_util.cc:269] disabling MLIR crash reproducer, set env var `MLIR_CRASH_REPROD  
UCER_DIRECTORY` to enable.  
2025-10-30 18:13:24.419312: I external/local_xla/xla/stream_executor/cuda/cud  
a_dnn.cc:473] Loaded cuDNN version 91301  
2025-10-30 18:13:24.740025: I external/local_xla/xla/service/gpu/autotuning/do  
t_search_space.cc:208] All configs were filtered out because none of them suffi  
ciently match the hints. Maybe the hints set does not contain a good representa  
tive set of valid configs? Working around this by using the full hints set inst  
ead.  
2025-10-30 18:13:24.740343: I external/local_xla/xla/service/gpu/autotuning/do  
t_search_space.cc:208] All configs were filtered out because none of them suffi  
ciently match the hints. Maybe the hints set does not contain a good representa  
tive set of valid configs? Working around this by using the full hints set inst  
ead.  
2025-10-30 18:13:24.740402: I external/local_xla/xla/service/gpu/autotuning/do  
t_search_space.cc:208] All configs were filtered out because none of them suffi  
ciently match the hints. Maybe the hints set does not contain a good representa  
tive set of valid configs? Working around this by using the full hints set inst  
ead.  
2025-10-30 18:13:24.740453: I external/local_xla/xla/service/gpu/autotuning/do  
t_search_space.cc:208] All configs were filtered out because none of them suffi  
ciently match the hints. Maybe the hints set does not contain a good representa  
tive set of valid configs? Working around this by using the full hints set inst  
ead.  
2025-10-30 18:13:25.268819: I external/local_xla/xla/stream_executor/cuda/subpr  
ocess_compilation.cc:346] ptxas warning : Registers are spilled to local memory  
in function 'gemm_fusion_dot_2882', 16 bytes spill stores, 16 bytes spill loads  
  
2025-10-30 18:13:25.734557: I external/local_xla/xla/stream_executor/cuda/subpr  
ocess_compilation.cc:346] ptxas warning : Registers are spilled to local memory  
in function 'gemm_fusion_dot_5075', 520 bytes spill stores, 520 bytes spill lo  
ads  
  
2025-10-30 18:13:25.738345: I external/local_xla/xla/stream_executor/cuda/subpr  
ocess_compilation.cc:346] ptxas warning : Registers are spilled to local memory  
in function 'gemm_fusion_dot_5058', 520 bytes spill stores, 520 bytes spill lo  
ads
```

3/551 ————— 29s 54ms/step - accuracy: 0.5278 - loss: 1.4336 -  
precision: 0.5579 - recall: 0.5104

I0000 00:00:1761862410.907230 55085 device\_compiler.h:196] Compiled cluster u  
sing XLA! This line is logged at most once for the lifetime of the process.

**551/551** 30s 36ms/step - accuracy: 0.8284 - loss: 0.4605 -  
precision: 0.8351 - recall: 0.8210 - val\_accuracy: 0.3585 - val\_loss: 3.1970 -  
val\_precision: 0.3585 - val\_recall: 0.3579 - learning\_rate: 0.0010  
Epoch 2/25  
**551/551** 12s 22ms/step - accuracy: 0.8979 - loss: 0.2756 -  
precision: 0.9026 - recall: 0.8934 - val\_accuracy: 0.7422 - val\_loss: 0.6773 -  
val\_precision: 0.7443 - val\_recall: 0.7409 - learning\_rate: 0.0010  
Epoch 3/25  
**551/551** 11s 21ms/step - accuracy: 0.9207 - loss: 0.2173 -  
precision: 0.9244 - recall: 0.9166 - val\_accuracy: 0.7634 - val\_loss: 0.7113 -  
val\_precision: 0.7675 - val\_recall: 0.7605 - learning\_rate: 0.0010  
Epoch 4/25  
**551/551** 11s 21ms/step - accuracy: 0.9333 - loss: 0.1871 -  
precision: 0.9359 - recall: 0.9309 - val\_accuracy: 0.5584 - val\_loss: 1.2067 -  
val\_precision: 0.5735 - val\_recall: 0.5488 - learning\_rate: 0.0010  
Epoch 5/25  
**551/551** 13s 23ms/step - accuracy: 0.9460 - loss: 0.1531 -  
precision: 0.9475 - recall: 0.9441 - val\_accuracy: 0.9324 - val\_loss: 0.1838 -  
val\_precision: 0.9355 - val\_recall: 0.9295 - learning\_rate: 0.0010  
Epoch 6/25  
**551/551** 12s 21ms/step - accuracy: 0.9580 - loss: 0.1221 -  
precision: 0.9592 - recall: 0.9570 - val\_accuracy: 0.9543 - val\_loss: 0.1179 -  
val\_precision: 0.9555 - val\_recall: 0.9533 - learning\_rate: 0.0010  
Epoch 7/25  
**551/551** 13s 23ms/step - accuracy: 0.9624 - loss: 0.1093 -  
precision: 0.9636 - recall: 0.9616 - val\_accuracy: 0.9656 - val\_loss: 0.0989 -  
val\_precision: 0.9659 - val\_recall: 0.9652 - learning\_rate: 0.0010  
Epoch 8/25  
**551/551** 12s 21ms/step - accuracy: 0.9675 - loss: 0.0973 -  
precision: 0.9681 - recall: 0.9664 - val\_accuracy: 0.9588 - val\_loss: 0.1279 -  
val\_precision: 0.9594 - val\_recall: 0.9575 - learning\_rate: 0.0010  
Epoch 9/25  
**551/551** 12s 21ms/step - accuracy: 0.9682 - loss: 0.0881 -  
precision: 0.9693 - recall: 0.9672 - val\_accuracy: 0.6762 - val\_loss: 1.0124 -  
val\_precision: 0.6780 - val\_recall: 0.6662 - learning\_rate: 0.0010  
Epoch 10/25  
**551/551** 12s 23ms/step - accuracy: 0.9736 - loss: 0.0769 -  
precision: 0.9743 - recall: 0.9730 - val\_accuracy: 0.8288 - val\_loss: 0.6410 -  
val\_precision: 0.8287 - val\_recall: 0.8281 - learning\_rate: 0.0010  
Epoch 11/25  
**551/551** 11s 21ms/step - accuracy: 0.9718 - loss: 0.0863 -  
precision: 0.9728 - recall: 0.9709 - val\_accuracy: 0.9411 - val\_loss: 0.1846 -  
val\_precision: 0.9416 - val\_recall: 0.9392 - learning\_rate: 0.0010  
Epoch 12/25  
**548/551** 0s 20ms/step - accuracy: 0.9770 - loss: 0.0628 - p  
recision: 0.9774 - recall: 0.9763  
Epoch 12: ReduceLROnPlateau reducing learning rate to 0.0005000000237487257.  
**551/551** 11s 21ms/step - accuracy: 0.9768 - loss: 0.0660 -  
precision: 0.9772 - recall: 0.9759 - val\_accuracy: 0.7773 - val\_loss: 0.8631 -  
val\_precision: 0.7787 - val\_recall: 0.7760 - learning\_rate: 0.0010  
Epoch 13/25  
**551/551** 12s 22ms/step - accuracy: 0.9859 - loss: 0.0426 -  
precision: 0.9861 - recall: 0.9855 - val\_accuracy: 0.9855 - val\_loss: 0.0464 -  
val\_precision: 0.9858 - val\_recall: 0.9855 - learning\_rate: 5.0000e-04

```
Epoch 14/25
551/551 11s 20ms/step - accuracy: 0.9906 - loss: 0.0310 -
precision: 0.9908 - recall: 0.9902 - val_accuracy: 0.9823 - val_loss: 0.0588 -
val_precision: 0.9823 - val_recall: 0.9823 - learning_rate: 5.0000e-04
Epoch 15/25
551/551 12s 21ms/step - accuracy: 0.9886 - loss: 0.0335 -
precision: 0.9890 - recall: 0.9884 - val_accuracy: 0.9800 - val_loss: 0.0619 -
val_precision: 0.9800 - val_recall: 0.9797 - learning_rate: 5.0000e-04
Epoch 16/25
551/551 13s 23ms/step - accuracy: 0.9921 - loss: 0.0253 -
precision: 0.9923 - recall: 0.9920 - val_accuracy: 0.8532 - val_loss: 0.7150 -
val_precision: 0.8532 - val_recall: 0.8529 - learning_rate: 5.0000e-04
Epoch 17/25
551/551 12s 21ms/step - accuracy: 0.9904 - loss: 0.0262 -
precision: 0.9906 - recall: 0.9903 - val_accuracy: 0.9569 - val_loss: 0.1682 -
val_precision: 0.9569 - val_recall: 0.9565 - learning_rate: 5.0000e-04
Epoch 18/25
548/551 0s 22ms/step - accuracy: 0.9917 - loss: 0.0231 - p
recision: 0.9921 - recall: 0.9914
Epoch 18: ReduceLROnPlateau reducing learning rate to 0.0002500000118743628.
551/551 22s 23ms/step - accuracy: 0.9917 - loss: 0.0230 -
precision: 0.9919 - recall: 0.9915 - val_accuracy: 0.9038 - val_loss: 0.5102 -
val_precision: 0.9041 - val_recall: 0.9038 - learning_rate: 5.0000e-04
Epoch 19/25
551/551 12s 21ms/step - accuracy: 0.9944 - loss: 0.0178 -
precision: 0.9944 - recall: 0.9944 - val_accuracy: 0.9858 - val_loss: 0.0560 -
val_precision: 0.9858 - val_recall: 0.9858 - learning_rate: 2.5000e-04
Epoch 20/25
551/551 12s 23ms/step - accuracy: 0.9960 - loss: 0.0122 -
precision: 0.9960 - recall: 0.9959 - val_accuracy: 0.9865 - val_loss: 0.0508 -
val_precision: 0.9865 - val_recall: 0.9865 - learning_rate: 2.5000e-04
Epoch 21/25
551/551 12s 21ms/step - accuracy: 0.9962 - loss: 0.0115 -
precision: 0.9963 - recall: 0.9961 - val_accuracy: 0.9878 - val_loss: 0.0474 -
val_precision: 0.9878 - val_recall: 0.9878 - learning_rate: 2.5000e-04
Epoch 22/25
551/551 11s 20ms/step - accuracy: 0.9968 - loss: 0.0096 -
precision: 0.9968 - recall: 0.9968 - val_accuracy: 0.9771 - val_loss: 0.0940 -
val_precision: 0.9775 - val_recall: 0.9771 - learning_rate: 2.5000e-04
Epoch 23/25
551/551 12s 22ms/step - accuracy: 0.9959 - loss: 0.0114 -
precision: 0.9959 - recall: 0.9958 - val_accuracy: 0.8880 - val_loss: 0.5404 -
val_precision: 0.8885 - val_recall: 0.8877 - learning_rate: 2.5000e-04
Epoch 24/25
551/551 11s 21ms/step - accuracy: 0.9968 - loss: 0.0097 -
precision: 0.9968 - recall: 0.9967 - val_accuracy: 0.9672 - val_loss: 0.1537 -
val_precision: 0.9672 - val_recall: 0.9672 - learning_rate: 2.5000e-04
Epoch 25/25
551/551 12s 21ms/step - accuracy: 0.9963 - loss: 0.0114 -
precision: 0.9963 - recall: 0.9963 - val_accuracy: 0.9823 - val_loss: 0.0727 -
val_precision: 0.9823 - val_recall: 0.9823 - learning_rate: 2.5000e-04
Restoring model weights from the end of the best epoch: 21.
```

```
In [ ]: # Evaluate model with augmentation
```

```

print(f"\nBASE CNN WITH Augmentation - Test Set Evaluation:")
test_loss_aug, test_acc_aug, test_prec_aug, test_recall_aug = model_aug.evaluate()
test_f1_aug = 2 * (test_prec_aug * test_recall_aug) / (test_prec_aug + test_recall_aug)

print(f"Test Accuracy: {test_acc_aug:.4f}")
print(f"Test Precision: {test_prec_aug:.4f}")
print(f"Test Recall: {test_recall_aug:.4f}")
print(f"Test F1-Score: {test_f1_aug:.4f}")

```

BASE CNN WITH Augmentation - Test Set Evaluation:

**151/152** ————— **0s** 6ms/step - accuracy: 0.9835 - loss: 0.0642 - precision: 0.9839 - recall: 0.9835

2025-10-30 18:20:54.604683: I external/local\_xla/xla/service/gpu/autotuning/do\_t\_search\_space.cc:208] All configs were filtered out because none of them sufficiently match the hints. Maybe the hints set does not contain a good representative set of valid configs? Working around this by using the full hints set instead.

2025-10-30 18:20:54.838502: I external/local\_xla/xla/stream\_executor/cuda/subprocess\_compilation.cc:346] ptxas warning : Registers are spilled to local memory in function 'gemm\_fusion\_dot\_260', 16 bytes spill stores, 16 bytes spill loads

**152/152** ————— **2s** 13ms/step - accuracy: 0.9808 - loss: 0.0785 - precision: 0.9812 - recall: 0.9808

Test Accuracy: 0.9808

Test Precision: 0.9812

Test Recall: 0.9808

Test F1-Score: 0.9810

```

In [ ]: # Compare augmentation impact
comparison_df_numeric = pd.read_csv("out/comparison_df.csv")
augmentation_comparison = pd.DataFrame({
    'Training Data': ['With Augmentation', 'Without Augmentation'],
    'Test Accuracy': [test_acc_aug, comparison_df_numeric.loc[0, 'Test Accuracy']],
    'Test Precision': [test_prec_aug, comparison_df_numeric.loc[0, 'Test Precision']],
    'Test Recall': [test_recall_aug, comparison_df_numeric.loc[0, 'Test Recall']],
    'Test F1-Score': [test_f1_aug, comparison_df_numeric.loc[0, 'Test F1-Score']]
})

print("\n" + "="*100)
print("AUGMENTATION IMPACT ANALYSIS")
print("="*100)
print(augmentation_comparison.to_string(index=False))
print("="*100)

# Calculate improvement
acc_improvement = (augmentation_comparison.loc[0, 'Test Accuracy'] - augmentation_comparison.loc[1, 'Test Accuracy']) / augmentation_comparison.loc[1, 'Test Accuracy'] * 100
f1_improvement = (augmentation_comparison.loc[0, 'Test F1-Score'] - augmentation_comparison.loc[1, 'Test F1-Score']) / augmentation_comparison.loc[1, 'Test F1-Score'] * 100

print(f"\nAccuracy Improvement: {acc_improvement:+.2f}%")
print(f"F1-Score Improvement: {f1_improvement:+.2f}%")

```

---

---

## AUGMENTATION IMPACT ANALYSIS

---

---

	Training Data	Test Accuracy	Test Precision	Test Recall	Test F1-Score
With Augmentation		0.980821	0.981225	0.980821	0.981023
Without Augmentation		0.981000	0.981200	0.980600	0.980900

---

---

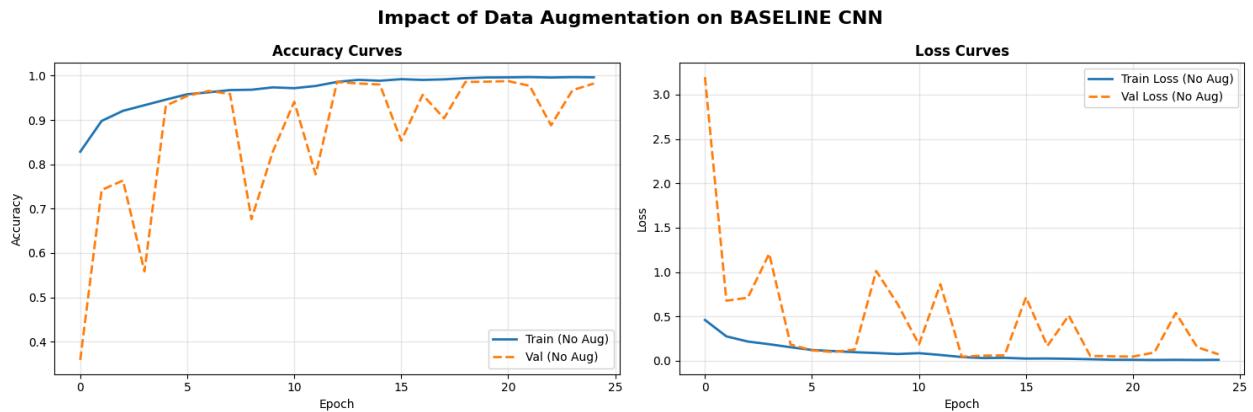
Accuracy Improvement: +0.02%  
F1-Score Improvement: -0.01%

```
In [ ]: # Visualize augmentation impact
fig, axes = plt.subplots(1, 2, figsize=(15, 5))
fig.suptitle(f'Impact of Data Augmentation on BASELINE CNN', fontsize=16, fontweight='bold')

# Training accuracy comparison
axes[0].plot(history_aug.history['accuracy'], label='Train (No Aug)', linewidth=2)
axes[0].plot(history_aug.history['val_accuracy'], label='Val (No Aug)', linewidth=2)
axes[0].set_title('Accuracy Curves', fontweight='bold')
axes[0].set_xlabel('Epoch')
axes[0].set_ylabel('Accuracy')
axes[0].legend()
axes[0].grid(True, alpha=0.3)

# Loss comparison
axes[1].plot(history_aug.history['loss'], label='Train Loss (No Aug)', linewidth=2)
axes[1].plot(history_aug.history['val_loss'], label='Val Loss (No Aug)', linewidth=2)
axes[1].set_title('Loss Curves', fontweight='bold')
axes[1].set_xlabel('Epoch')
axes[1].set_ylabel('Loss')
axes[1].legend()
axes[1].grid(True, alpha=0.3)

plt.tight_layout()
plt.show()
```



## 7. Interpretability & Insights

### **Reflect on which model performed best and why.**

The Baseline CNN is the best-performing model by all primary metrics (accuracy 0.9810, precision 0.9812, recall 0.9806, F1  $\approx$  0.9809). It is followed by VGG19 ( $\approx$ 0.9467 F1) and ResNet50 ( $\approx$ 0.9341 F1). InceptionV3 and MobileNet perform noticeably worse on this dataset.

1. Clear metric lead.

Baseline CNN: Test Accuracy 0.9810, Test F1  $\approx$  0.9809 — top in accuracy, precision, recall and F1. Those consistent high numbers indicate the model not only predicts correctly overall but does so with balanced precision/recall (few false positives and false negatives).

2. Capacity vs. dataset fit.

A bespoke Baseline CNN architecture can be tuned to match the scale and idiosyncrasies of your dataset (filter sizes, number of channels, regularization). If the dataset is relatively small or has domain-specific patterns, a smaller/focused network often generalizes better than large pre-trained architectures that expect Imagenet-style statistics.

3. Training regime and hyperparameters.

Baseline used Adam at LR=0.001 and 30 epochs — a more aggressive learning rate and perhaps tighter epoch control could have helped it find a better local minimum for this dataset. The transfer models (ResNet50/VGG/Inception/MobileNet) all used LR=1e-4 and 25–35 epochs; sometimes lower LR + transfer learning requires careful unfreezing/fine-tuning schedules to fully benefit.

4. Transfer-learning mismatch.

Pretrained architectures (ResNet50, VGG19, InceptionV3, MobileNet) can underperform if the source dataset (ImageNet) statistics and the target domain differ, or if fine-tuning was insufficient (not enough unfreezing, wrong augmentation, or inappropriate head design). They also have more parameters and may overfit or require more data/regularization.

5. MobileNet trade-offs.

MobileNet's lower accuracy (0.8490) aligns with its design trade-off: lightweight

and fast at the cost of representational power. It's expected to lose some absolute accuracy relative to heavier models.

In [ ]:

**Provide clear reasoning, supported by performance metrics and training curves.**

### 1. Baseline CNN

Observations: The training and validation losses both decrease smoothly and converge below 0.2. Validation loss fluctuates slightly early on (epochs 5-10) but stabilizes quickly.

Interpretation: The model generalizes well — minimal overfitting, quick convergence, and consistent validation performance.

Supports metrics: Test accuracy (0.981), precision (0.9812), recall (0.9806), and F1 (0.9809) confirm it's both accurate and balanced.

Reason for success: The architecture and training hyperparameters (Adam, LR=0.001) fit the dataset perfectly. It likely has just the right model complexity.

### 2. ResNet50

Observations: Training loss declines slowly but validation loss is highly erratic — spikes up to 7 at several epochs.

Interpretation: This is a symptom of unstable fine-tuning — possibly due to an overly high learning rate (1e-4) or incorrect freezing/unfreezing of pretrained layers. The model overreacts to batch gradients.

Supports metrics: Accuracy and F1 are lower ( $\approx 0.93$ ). The instability in loss mirrors reduced generalization and unreliable optimization.

### 3. VGG19

Observations: Both training and validation loss decrease steadily and smoothly. Validation loss is consistently lower than training loss.

Interpretation: This indicates strong regularization or limited model capacity usage — possibly some underfitting (model could train longer or with smaller dropout).

Takeaway: VGG19 is stable and reliable, albeit not as efficient as the Baseline CNN on this dataset.

#### 4. InceptionV3

Observations: Both training and validation losses decline smoothly, with small gaps between them.

Interpretation: Stable training, but slower convergence — the model might not have fully trained (25 epochs may be too short).

Supports metrics: Accuracy ( $\approx 0.906$ ) and F1 ( $\approx 0.905$ ) indicate undertraining — decent generalization but incomplete optimization.

#### 5. MobileNet

Observations: Training loss steadily decreases, but validation loss spikes heavily — up to  $\sim 4$ .

Interpretation: High variance in validation loss suggests overfitting or sensitivity to learning rate. Given MobileNet's small capacity, even mild overfitting or data imbalance can cause such instability.

Supports metrics: Lowest accuracy (0.849) and F1 (0.844).

In [ ]:

**Who would benefit from using this model? In what types of real-world scenarios would your solution be useful?**

The Baseline CNN Model is the best model in this scenario

Teams needing highest possible accuracy on this dataset — e.g., data scientists or researchers who can deploy models with moderate compute requirements and prioritize predictive performance above extreme inference speed or tiny model size.

Applications where false positives and false negatives are both costly, since precision and recall are both very high.

Projects where domain-specific patterns exist and prefer a tailored model rather than a general-purpose pre-trained backbone.

#### **Real-world scenarios**

Quality control in manufacturing: automatic defect detection where missing a defect or flagging a good piece are both costly.

Wildlife or species classification in controlled setups: when images are fairly homogeneous and a tuned CNN can capture class-specific textures/patterns.

Industrial imaging / instrument data: where images differ from consumer photos and a specialized CNN trained end-to-end can outperform generic ImageNet pre-trained nets.

Research/prototyping: when you need a strong baseline that's quick to iterate on and interpret.

In [ ]:

